



# **Training Neural Networks on Embedded Devices Targeting Embedded Environments**

---

Prasanth Shaji, Deepak Venkataram

Master's Thesis  
Uppsala University



# Outline

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

## 1 Introduction

## 2 Embedded Linux

## 3 Neural Networks

## 4 Benchmark Applications

## 5 Discussion



# Embedded Devices



Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion



UPPSALA  
UNIVERSITET

# Neural Network Applications on Embedded Devices

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion





# Goal : Neural Network Training on Scania ECU



Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion



# Goal : Neural Network Training on Scania ECU

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

## Why Training on the ECU?

- Training on devices reduces network usage, ensures data privacy, etc.
- Repurpose existing ECUs to perform neural network related tasks

*Scope Repurpose Scania ECU, train neural network model, assess training performance*



# Goal : Neural Network Training on Scania ECU

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

## Why Training on the ECU?

- Training on devices reduces network usage, ensures data privacy, etc.
- Repurpose existing ECUs to perform neural network related tasks

**Scope** Repurpose Scania ECU, train neural network model, assess training performance



# Goal : Neural Network Training on Scania ECU

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

## Why Training on the ECU?

- Training on devices reduces network usage, ensures data privacy, etc.
- Repurpose existing ECUs to perform neural network related tasks

*Scope Repurpose Scania ECU, train neural network model, assess training performance*



# Goal : Neural Network Training on Scania ECU

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

## Why Training on the ECU?

- Training on devices reduces network usage, ensures data privacy, etc.
- Repurpose existing ECUs to perform neural network related tasks

**Scope** *Repurpose Scania ECU, train neural network model, assess training performance*



# Embedded Devices : Hardware Outline

## ARM | Silicon Vendor | System Maker

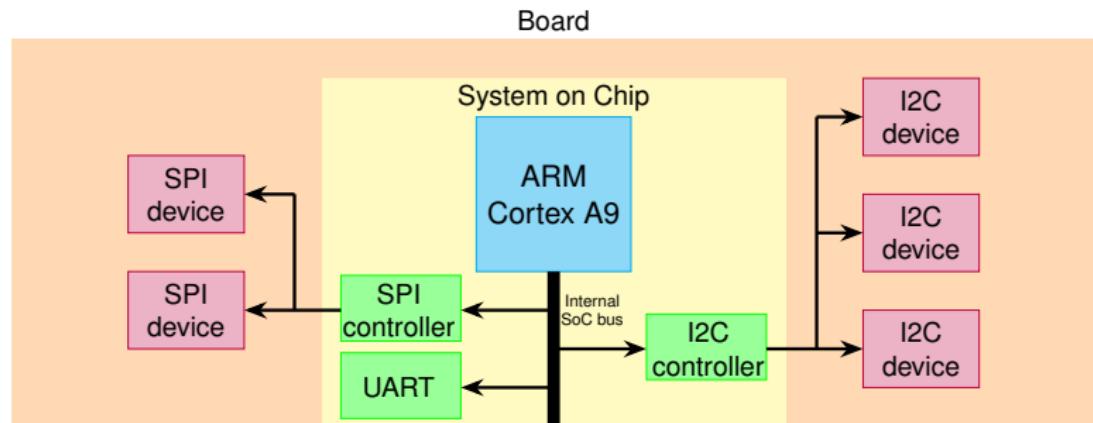
Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion





# ARM Boards : Examples

## Introduction

Embedded Linux

Neural Networks

Benchmark Applications

Discussion

SoC	Application Domain	Architecture	Processor Core
nRF51822	Ultra low power, BLE	ARM v6-M	ARM Cortex M0
AM2732	Automotive	ARM v7-R	ARM Cortex R5
AM3358	Industrial / IoT	ARM v7-A	ARM Cortex A8
i.MX6S	Multimedia applications	ARM v7-A	ARM Cortex A9
BCM2711	Educational / IoT	ARM v8-A	ARM Cortex-A72
Snapdragon 460	Smartphones	ARM v8-A	ARM Cortex-A73, A53
Google Tensor	Pixel 6	ARM v8.2-A	ARM Cortex-X1, A76, A55
Apple A16 Bionic	iPhone 14 Pro	ARM v8.6-A	APL1W10
Apple M1	Macbooks / Mac Mini / iPads	ARM v8.5-A	APL1102



# Embedded Linux

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

Operating System for Embedded Devices  
- Customized for Specific Hardware

Kernel, Bootloaders, Device Trees  
Cross Compiling Toolchains, Root Filesystems



# Embedded Linux

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

Operating System for Embedded Devices  
- Customized for Specific Hardware

Kernel, Bootloaders, Device Trees  
Cross Compiling Toolchains, Root Filesystems



# Build Systems

Introduction

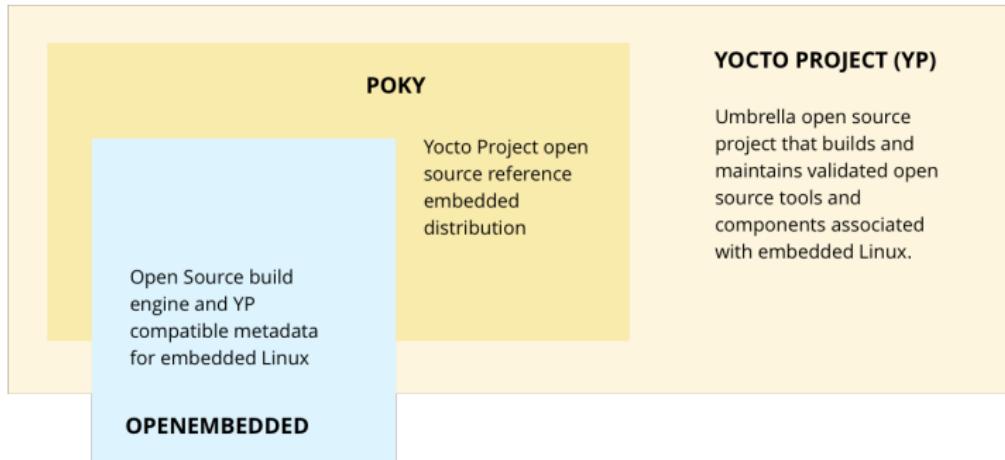
Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

## Yocto Project





# Development Setup

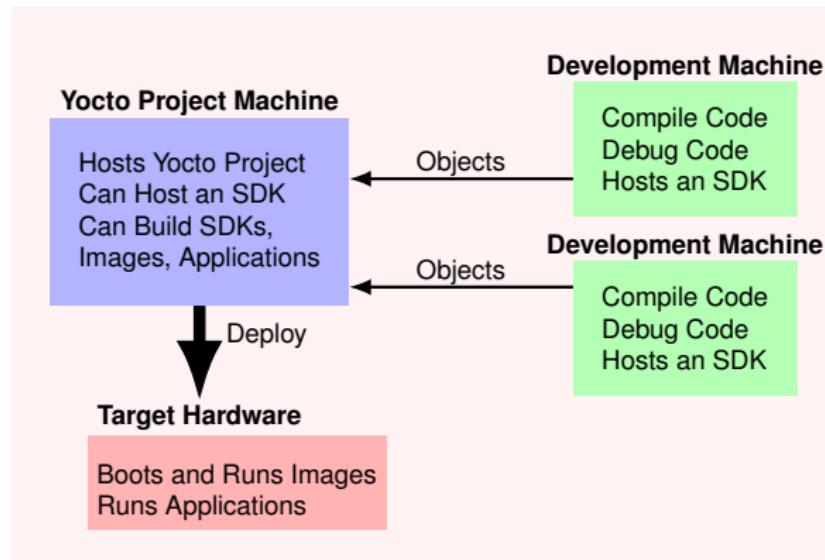
Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion





# Software Versions

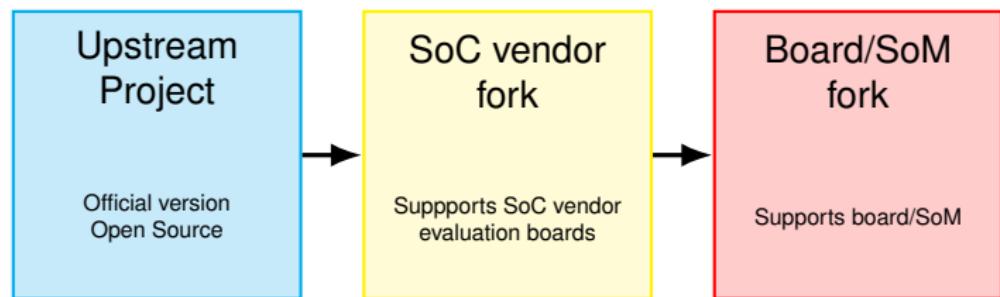
Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion





# Scania ECU to MCIMX6Q-SDB

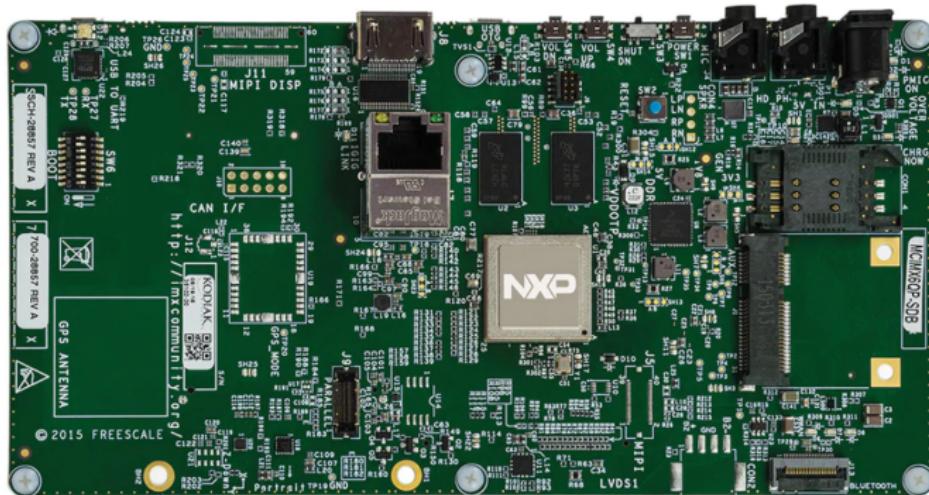
Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion





# Handwritten Digit Recognition

## Neural Network

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion



MNIST



# Neural Network

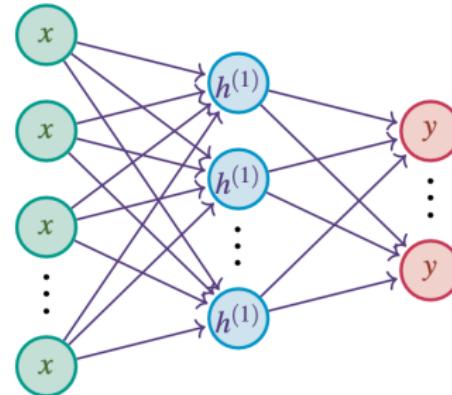
Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion



$$y_k = \sigma \left( \sum_{j=0}^m w_{kj} x_j \right)$$



# Neural Network Training

Introduction

Embedded  
Linux

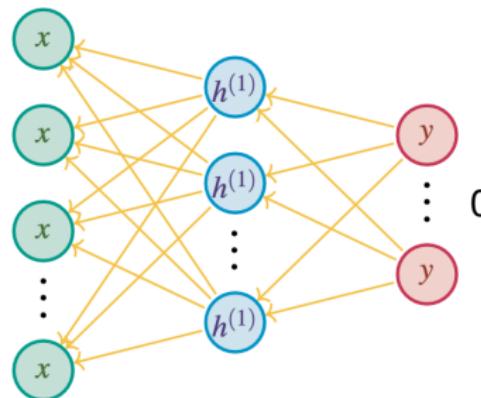
Neural  
Networks

Benchmark  
Applications

Discussion



label = 0





# Learning Algorithm

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

---

**Algorithm 1** Mini Batch Gradient Descent with learning rate  $\eta$  and the Mean Squared Error (MSE) cost function

---

**Require:** initial weights  $w^{(0)}$ , number of epochs  $E$ , batch size  $B$ , training data with  $T$  entries

**Ensure:** final weights  $w^{(E*T)}$

```
for  $e = 0 \rightarrow E - 1$  do
    for  $b = 0 \rightarrow T/B$  do
        for  $t = b * B \rightarrow (b + 1) * B$  do
            estimate  $\nabla \mathcal{L}(w^{(t)})$ 
            compute  $\Delta w^{(b)} += -\nabla \mathcal{L}(w^{(t)})$                                  $\triangleright \mathcal{L}$  here is MSE
        end for
         $w^{(e+1)} := w^{(e)} + \eta \Delta w^{(e)}$ 
    end for
end for
return  $w^{(T)}$ 
```

---



# Neural Network Inference

Introduction

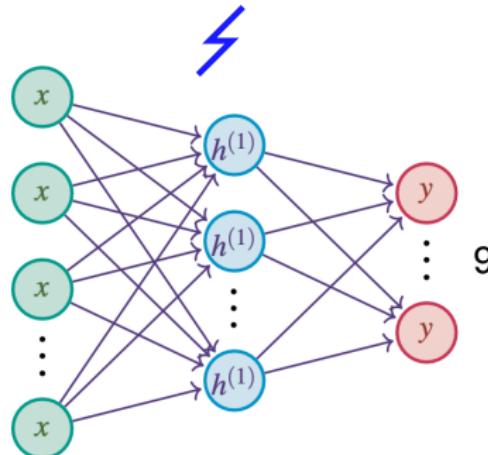
Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

9





# Traditional Paradigm

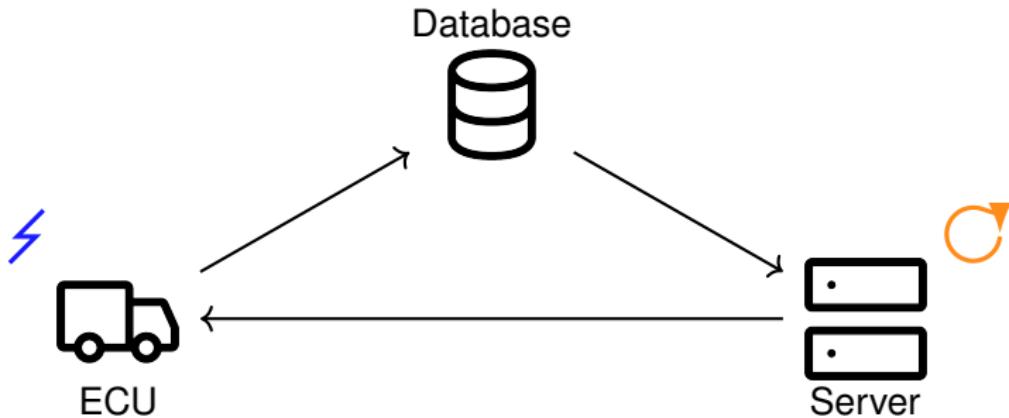
Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion





# Federated Learning

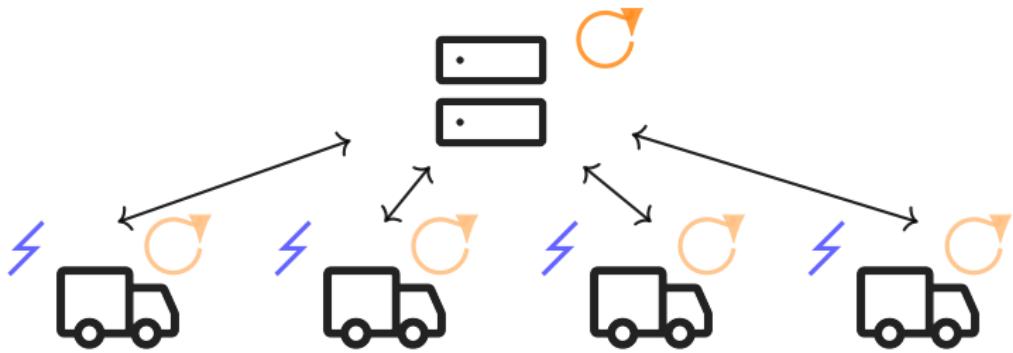
Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion





# HDR-NN Implementations

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

[github.com/paperbound/hdrnn-arm](https://github.com/paperbound/hdrnn-arm)

- C based HDR-NN
- C++ / Eigen
- C++ / libtorch (PyTorch)
- Python / Numpy



# C based HDR-NN

Introduction

Embedded  
LinuxNeural  
NetworksBenchmark  
Applications

Discussion

```
/*-HDR-Neural-Network*/
typedef struct
{
    float *bias;
    float *weights;
    float *nabla_w;
} Neuron;

typedef struct LayerT
{
    int size;
    int incidents;
    Neuron *neurons;
    float *activations;
    float *z_values;
    float *nabla_b;
    struct LayerT *next;
    struct LayerT *previous;
} Layer; // Network.layers except for input

typedef struct
{
    Layer *layers;
    int depth;
} Network; // HDRNN
```

```
static void feed_forward(Network *network, float *image)
{
    // First layer is the image
    float *activations = image;

    Layer *layer = network->layers;
    while (layer != NULL)
    {
        for (int i = 0; i < layer->size; i++)
        {
            Neuron *neuron = &layer->neurons[i];
            float zvalue = 0;
            for (int j = 0; j < layer->incidents; j++)
            {
                zvalue += neuron->weights[j] * activations[j];
            }
            layer->z_values[i] = zvalue + neuron->bias;
            layer->activations[i] = sigmoid(layer->z_values[i]);
        }
        activations = layer->activations;
        layer = layer->next;
    }
}
```



# C++ / Eigen based HDR-NN

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

- Eigen is a C++ template library for linear algebra
- Analogous to Python / Numpy



# C++ / libtorch based HDR-NN

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

- PyTorch doesn't support MCIMX6Q
- Porting libtorch using QEMU



# HDR-NN configurability

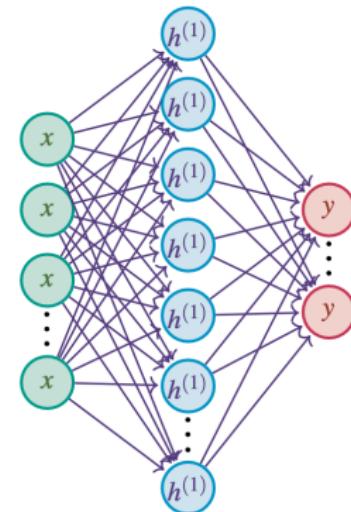
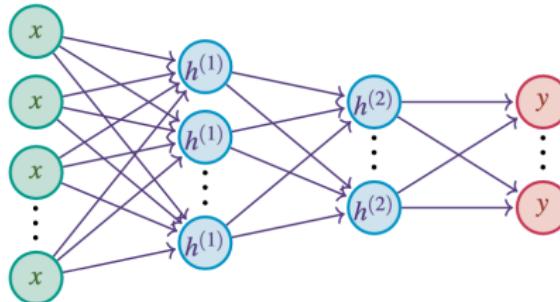
Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion





# HDR-NN configurability

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

```
→ ./bin/hdrnn
nothing to do
Usage: ./bin/hdrnn

      <command> [<args>]

Commands:
  infer [-i, --image IMAGE_PATH] [-n, --net c-math.nn]
  train [-s, --shape 32] [-e, --epochs 30]
        [-q, --quiet] [-bs, --batch_size=10]
        [-lr, --learning_rate 3] [-n, --net c-math.nn]
```



# Experiment Considerations

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

- Try multiple shapes m
- Compare Accuracies
- Compare Execution Time

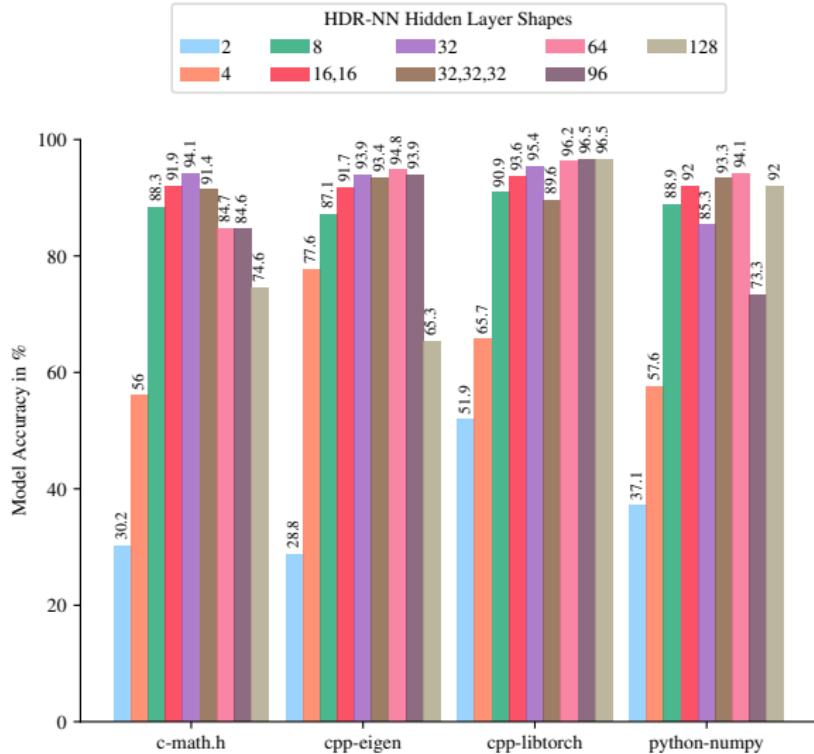


# Accuracy

Introduction

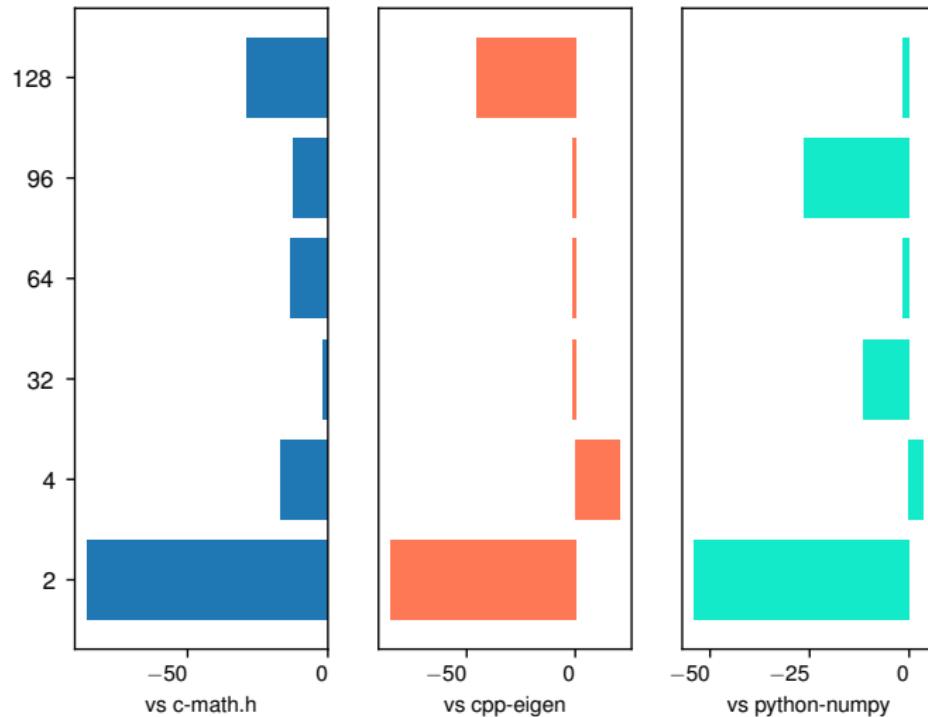
Embedded  
LinuxNeural  
NetworksBenchmark  
Applications

Discussion



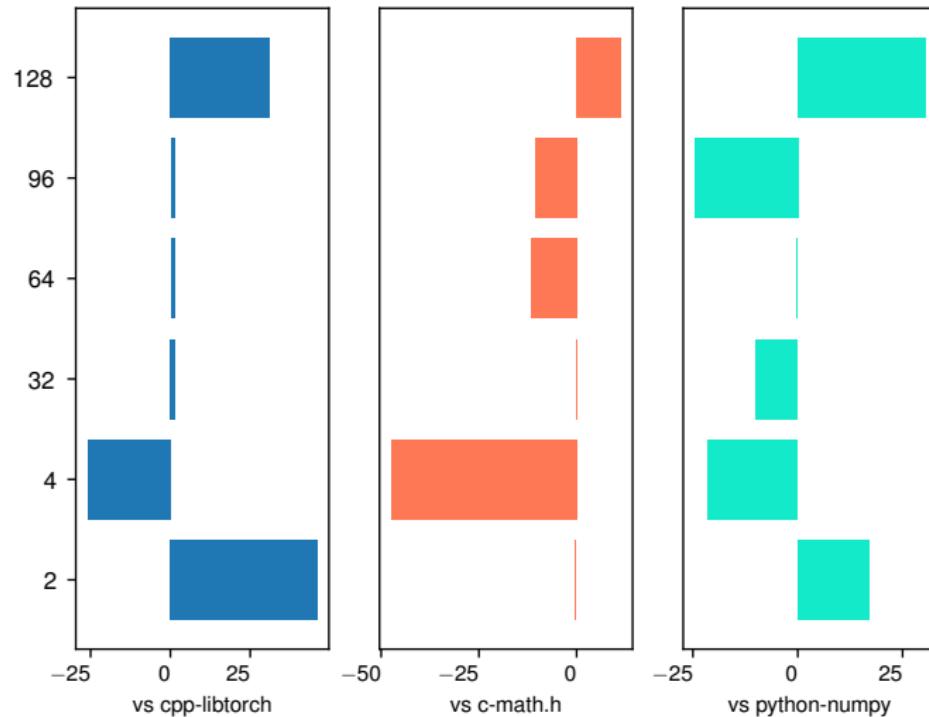


## Comparing Accuracies of cpp-libtorch



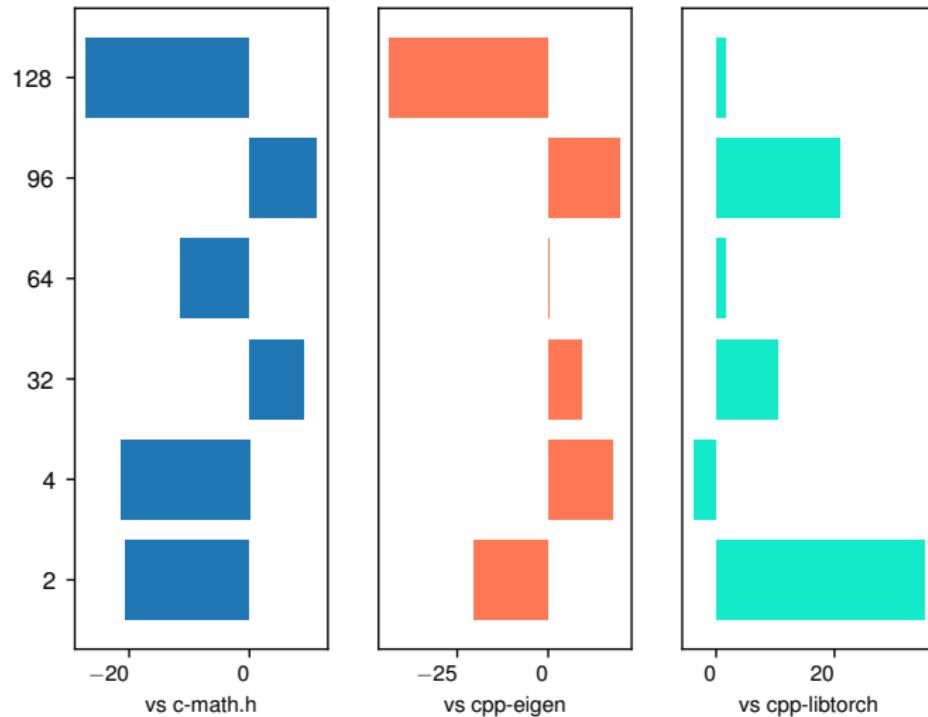


## Comparing Accuracies of cpp-eigen



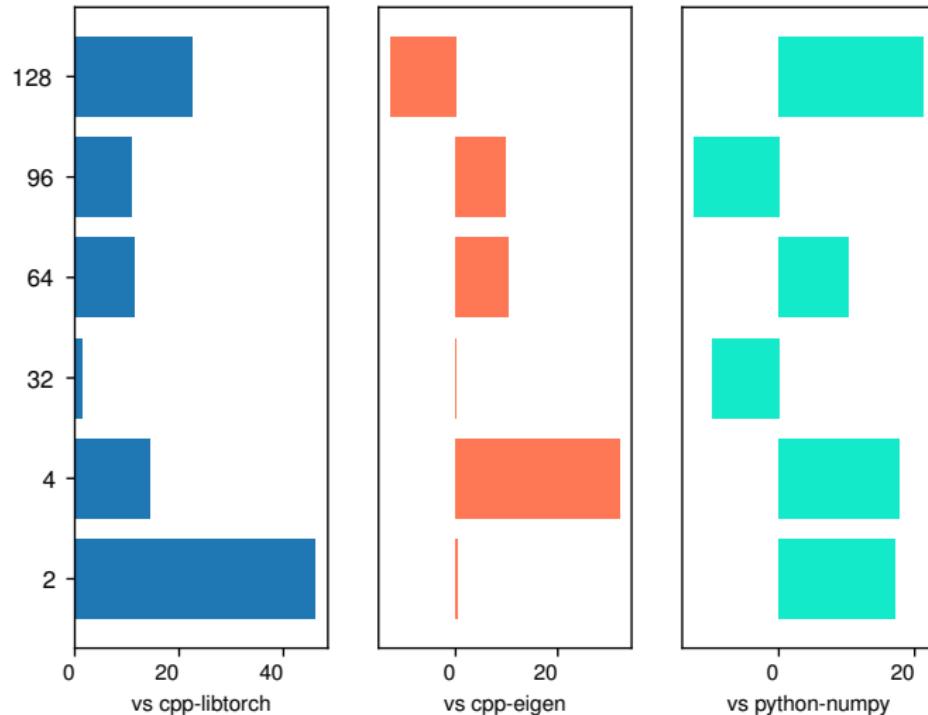


## Comparing Accuracies of python-numpy





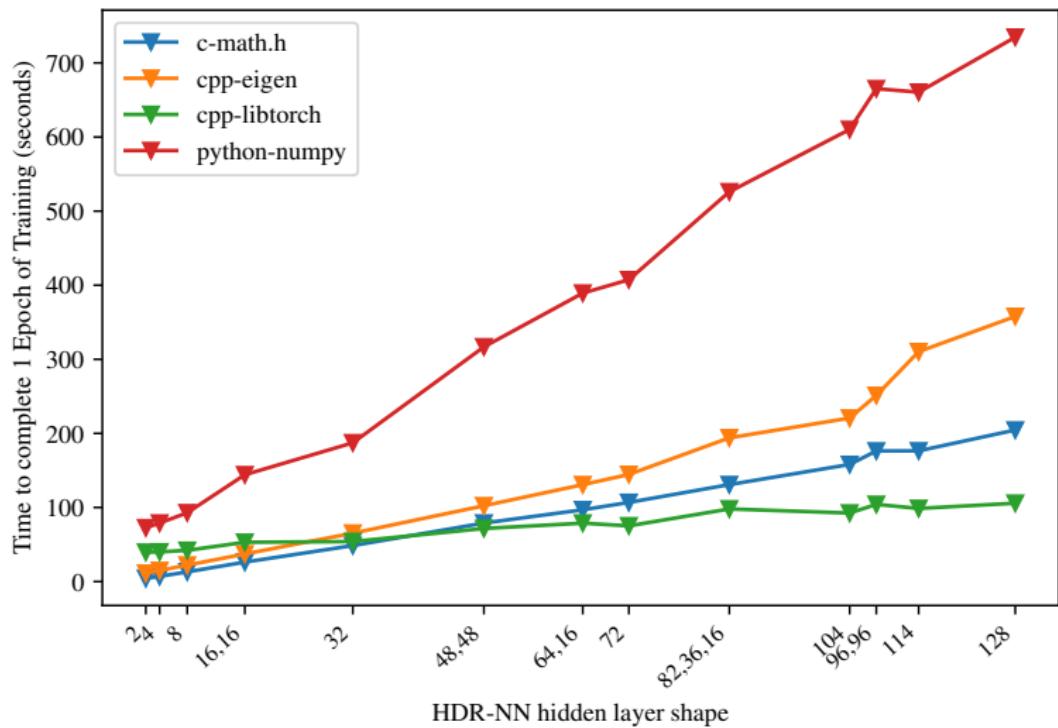
## Comparing Accuracies of c-math.h



# Training Time

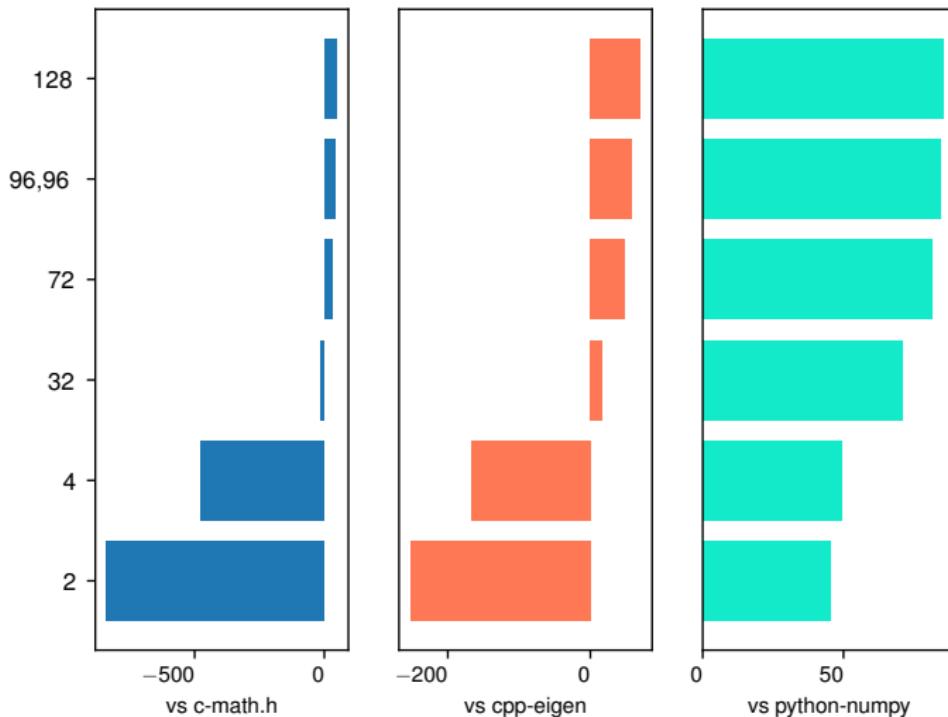
Introduction  
Embedded Linux  
Neural Networks  
Benchmark Applications

Discussion



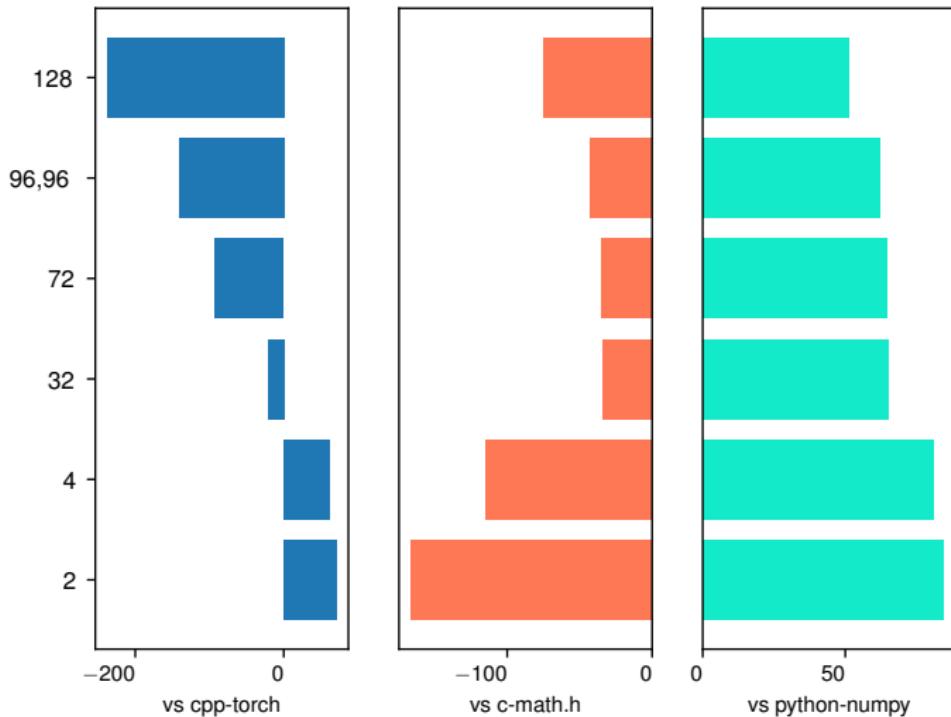


## Comparing Execution Times of cpp-libtorch



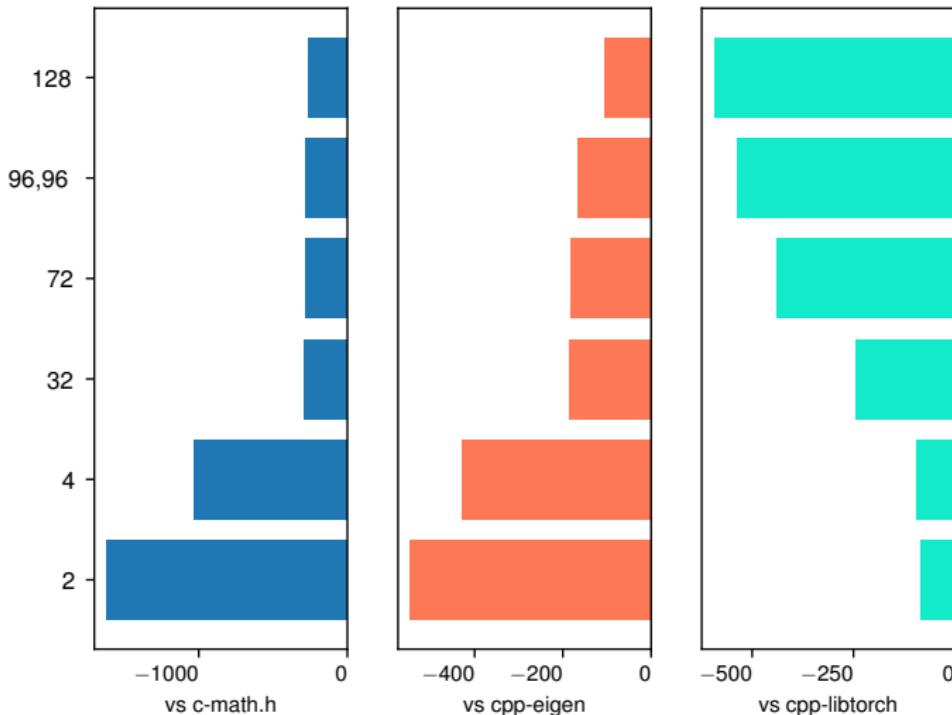


## Comparing Execution Times of cpp-eigen



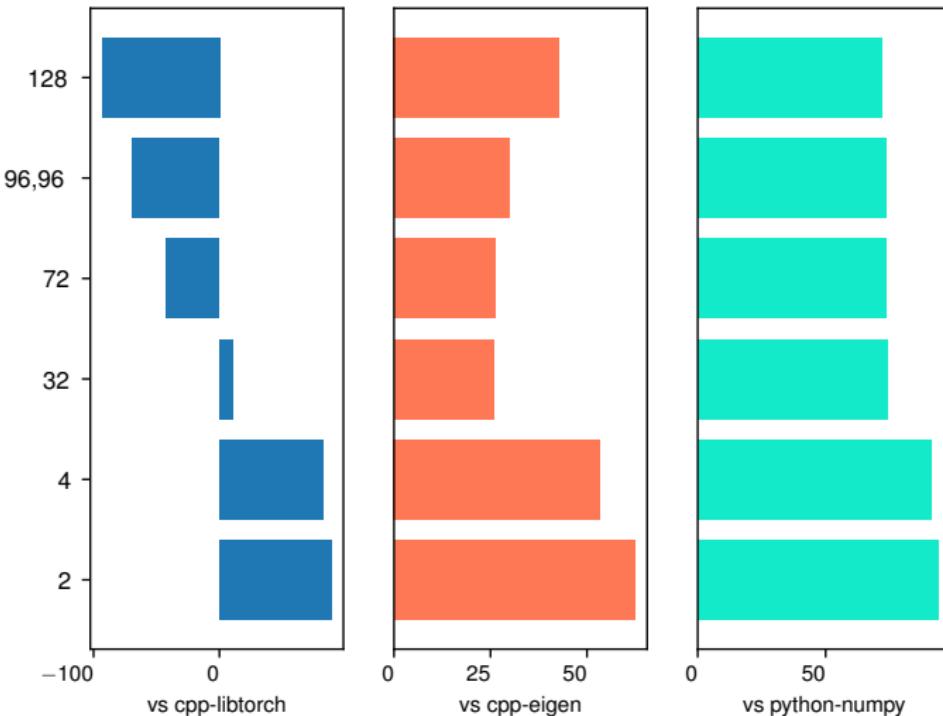


## Comparing Execution Times of python-numpy





## Comparing Execution Times of c-math.h





# Reverse Engineering Scania ECU

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

- Goal : Get a vanilla U-Boot running on ECU
- Extract device tree information
- Silent bootloader, No visibility into Board/SoM fork



# Neural Network Software Development on Embedded Devices

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

- Embedded Platforms are fragmented in both *hardware* and *software*
- C is easier to target harder to Implement
- Python commonly found in Neural Network Research, harder to port



# Future Work

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

- Reverse Engineering Efforts on the ECU
- Performance engineering the C based implementation
- Adding more target devices
- Adding more implementations (Tensorflow, TTE)



End

---

Introduction

Embedded  
Linux

Neural  
Networks

Benchmark  
Applications

Discussion

Thank you