



Training Neural Networks on Embedded Devices Targeting Embedded Environments

Prasanth Shaji, Deepak Venkataram

Master's Thesis
Uppsala University

August 20, 2023



NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion

Outline

1 NN in Embedded Devices

2 Embedded Linux

3 Neural Networks

4 Benchmark Applications

5 Discussion

6 Conclusion



Neural Network Applications on Embedded Devices

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion





Embedded Devices

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion





Goal : Neural Network training on Scania C300

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion





Goal : Neural Network training on Scania C300

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the C300?

- Repurpose existing C300 units to perform neural network related tasks
- Training on devices to reduce network bandwidth usage, ensure data privacy, etc.

Scope Repurpose Scania C300, train neural network model, assess training performance



Goal : Neural Network training on Scania C300

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the C300?

- Repurpose existing C300 units to perform neural network related tasks
- Training on devices to reduce network bandwidth usage, ensure data privacy, etc.

Scope Repurpose Scania C300, train neural network model, assess training performance



Goal : Neural Network training on Scania C300

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the C300?

- Repurpose existing C300 units to perform neural network related tasks
- Training on devices to reduce network bandwidth usage, ensure data privacy, etc.

Scope Repurpose Scania C300, train neural network model, assess training performance



ARM Boards

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion

SoC	Application Domain	Architecture	Processor Core
nRF51822	Ultra low power, BLE	ARM v6-M	ARM Cortex M0
AM2732	Automotive	ARM v7-R	ARM Cortex R5
AM3358	Industrial / IoT	ARM v7-A	ARM Cortex A8
i.MX6S	Multimedia applications	ARM v7-A	ARM Cortex A9
Kryo 240	Smartphones	ARM v8-A	ARM Cortex-A73, A53



ARM Boards

NN in
Embedded
Devices

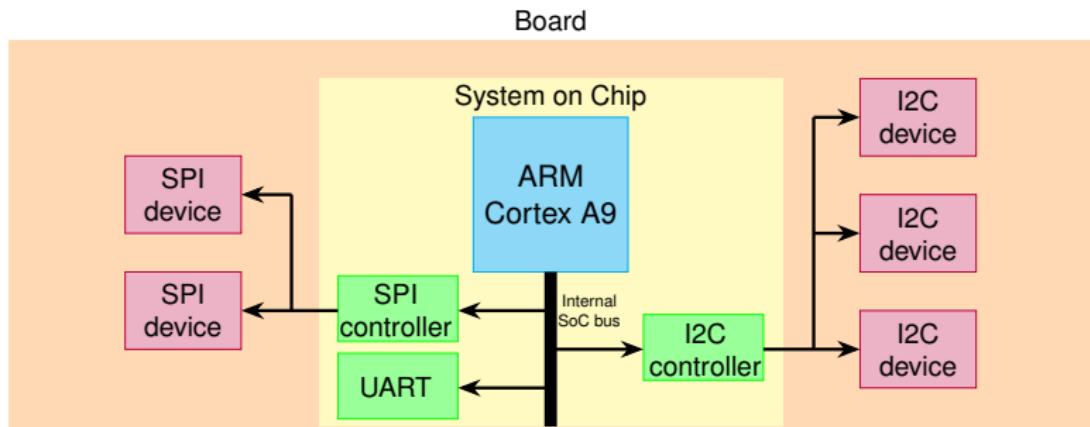
Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion





Embedded Linux

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion

Operating System for Embedded Devices
- Customized for Specific Hardware

Kernel, Bootloaders, Device Trees
Cross Compilation, Toolchains, Root Filesystems



Software Versions

NN in
Embedded
Devices

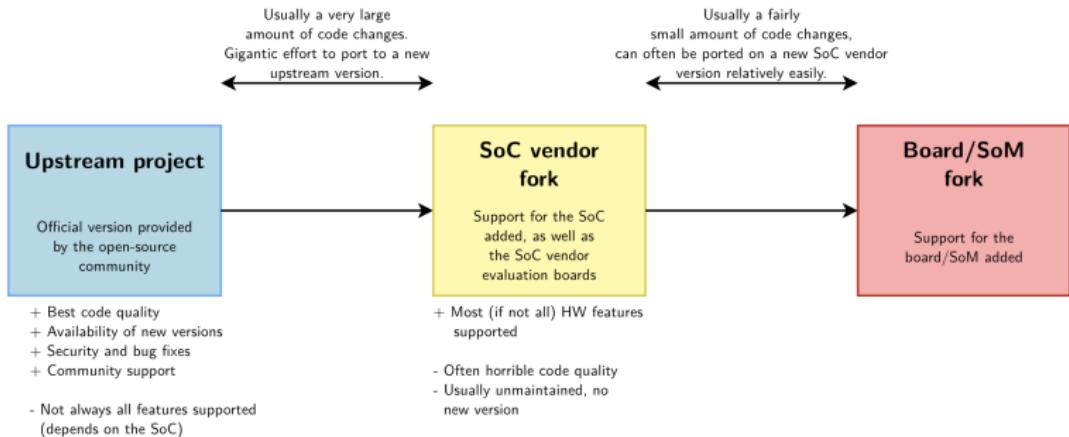
Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion





Build Systems

NN in
Embedded
Devices

Embedded
Linux

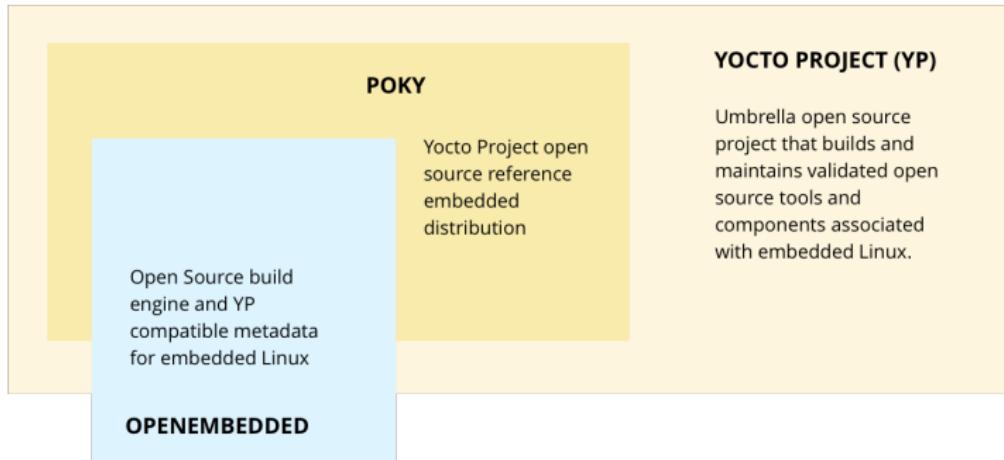
Neural
Networks

Benchmark
Applications

Discussion

Conclusion

Yocto Project





Yocto Project

NN in
Embedded
Devices

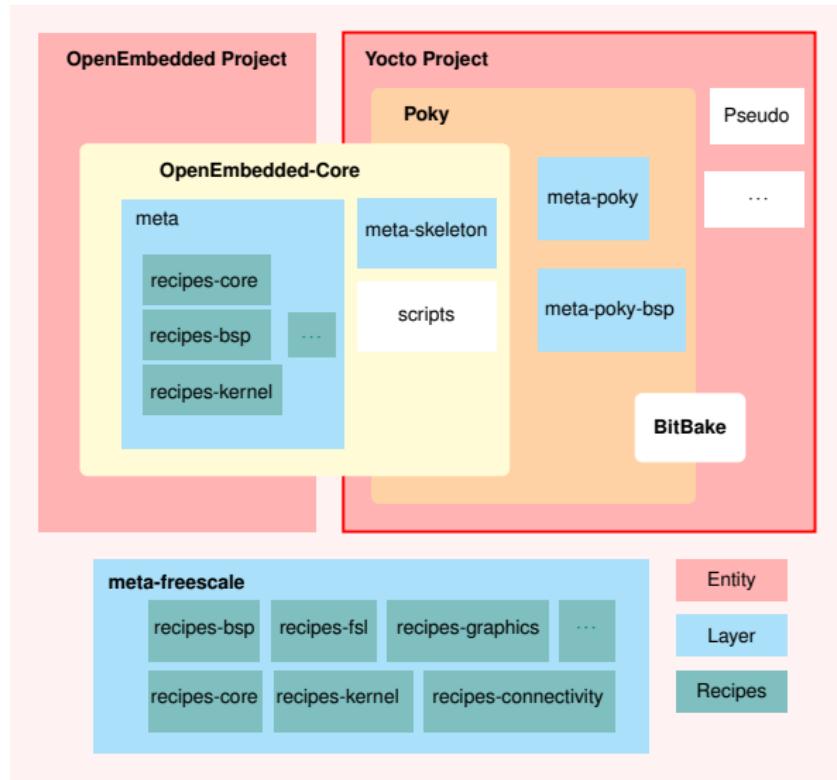
Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion





Yocto Project

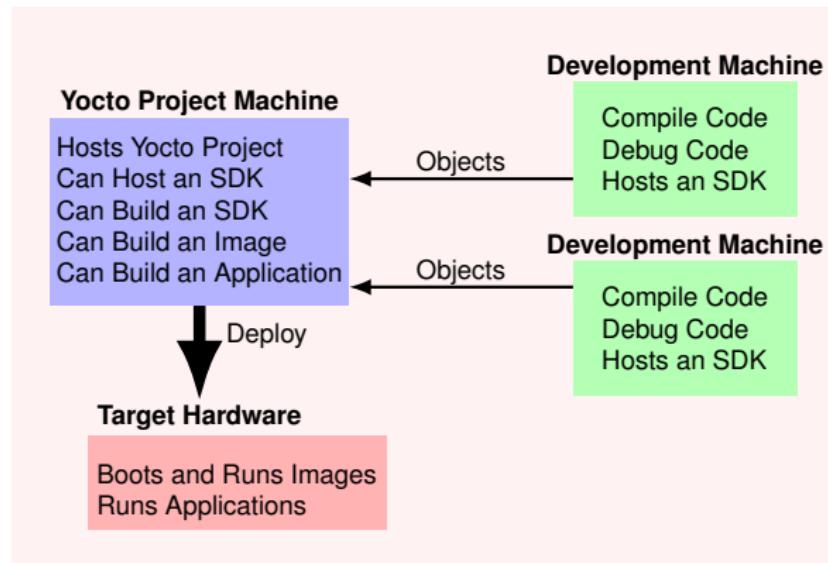
NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion





Scania C300 to MCIMX6Q-SDB

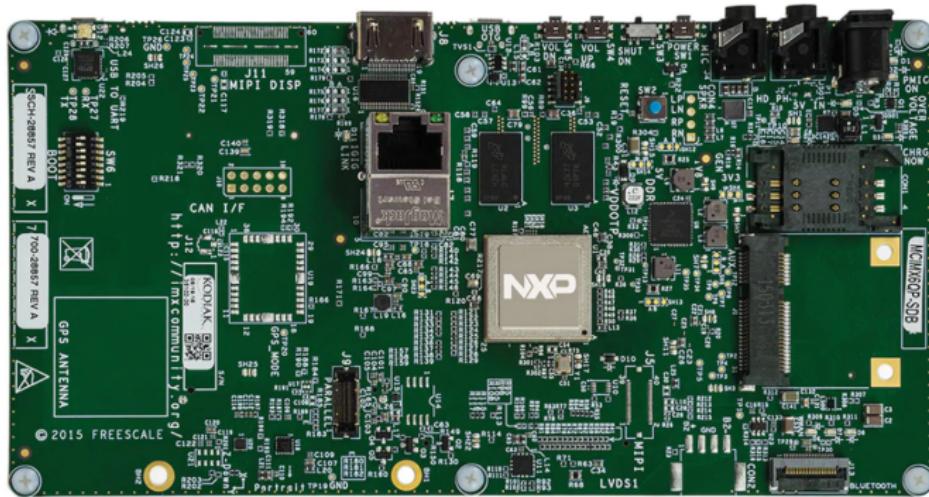
NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion





Training and Inference

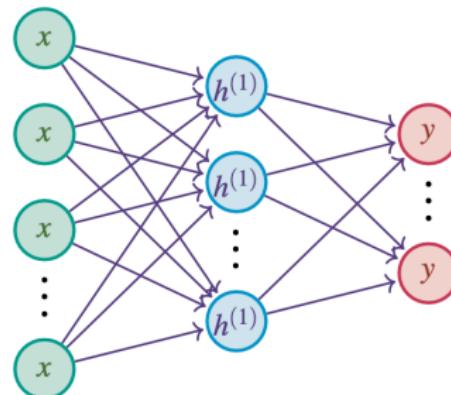
NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion





Traditional Paradigm

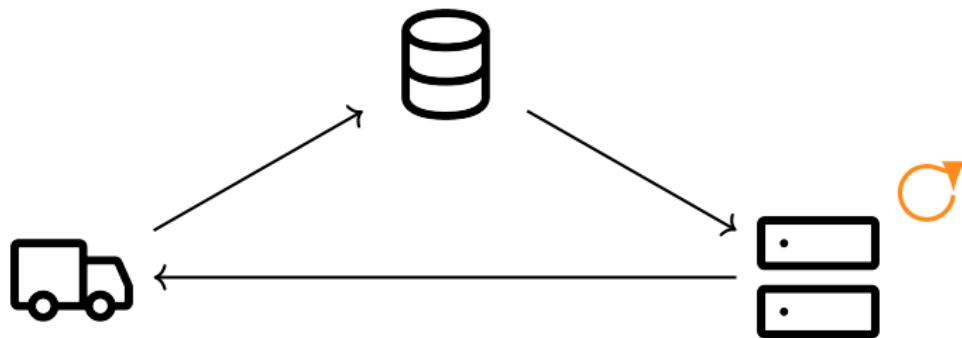
NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion





Federated Learning

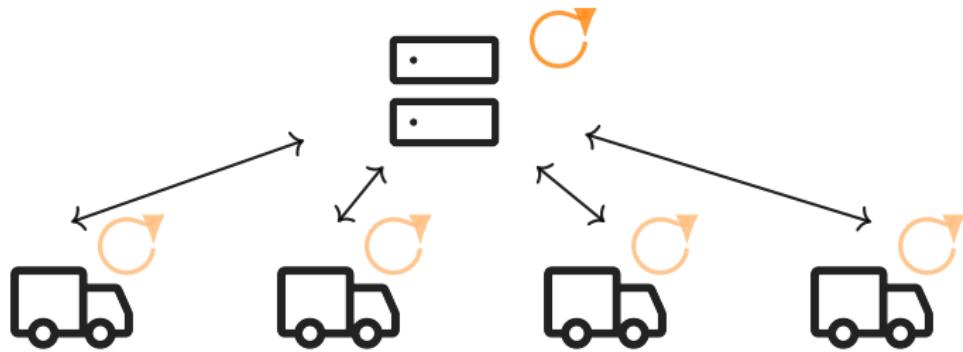
NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion





Handwritten Digit Recognition Neural Network

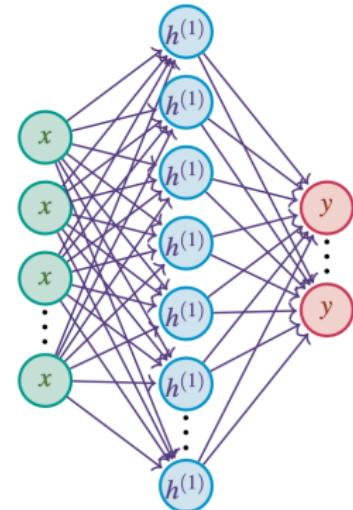
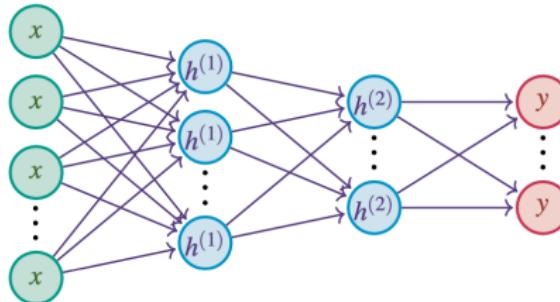
NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion





Learning Algorithm

Algorithm 1 Mini Batch Gradient Descent with learning rate η and the Mean Squared Error (MSE) cost function

Require: initial weights $w^{(0)}$, number of epochs E , batch size B , training data with T entries

Ensure: final weights $w^{(E*T)}$

```
for  $e = 0 \rightarrow E - 1$  do
    for  $b = 0 \rightarrow T/B$  do
        for  $t = b * B \rightarrow (b + 1) * B$  do
            estimate  $\nabla \mathcal{L}(w^{(t)})$ 
            compute  $\Delta w^{(b)} += -\nabla \mathcal{L}(w^{(t)})$  ▷  $\mathcal{L}$  here is MSE
        end for
         $w^{(e+1)} := w^{(e)} + \eta \Delta w^{(e)}$ 
    end for
end for
return  $w^{(T)}$ 
```



HDR-NN configurability

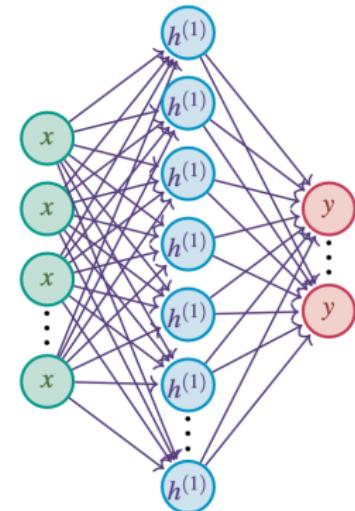
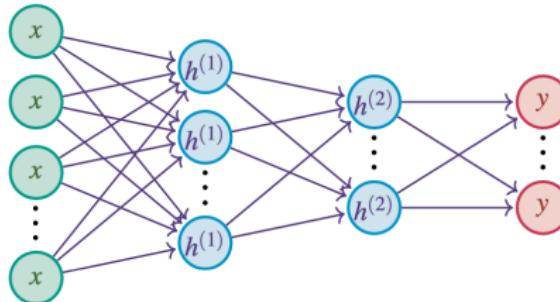
NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion
Conclusion





HDR-NN configurability

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion

```
→ ./bin/hdrnn
nothing to do
Usage: ./bin/hdrnn

      <command> [<args>]

Commands:
  infer [-i, --image IMAGE_PATH] [-n, --net c-math.nn]
  train [-s, --shape 32] [-e, --epochs 30]
        [-q, --quiet] [-bs, --batch_size=10]
        [-lr, --learning_rate 3] [-n, --net c-math.nn]
```



C based HDR-NN

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion

```
/*HDR.Neural.Network*/
typedef struct
{
    float *bias;
    float **weights;
    float **nabla_w;
} Neuron;

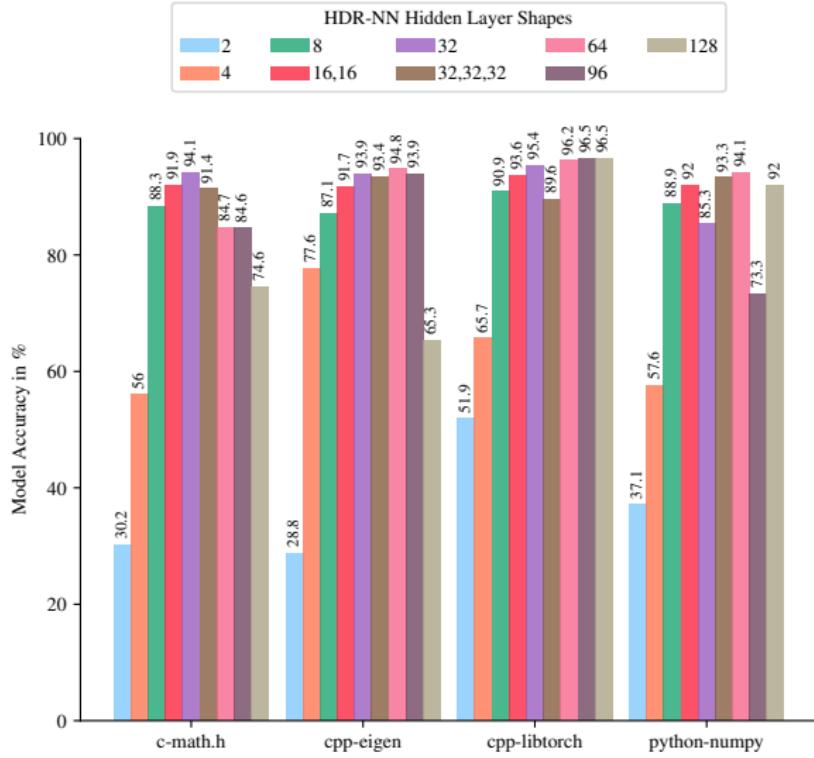
typedef struct LayerT
{
    int size;
    int incidents;
    Neuron **neurons;
    float **activations;
    float **z_values;
    float **nabla_b;
    struct LayerT *next;
    struct LayerT *previous;
} Layer; // Network.layers except for input

typedef struct
{
    Layer **layers;
    int depth;
} Network; // HDRNN
```

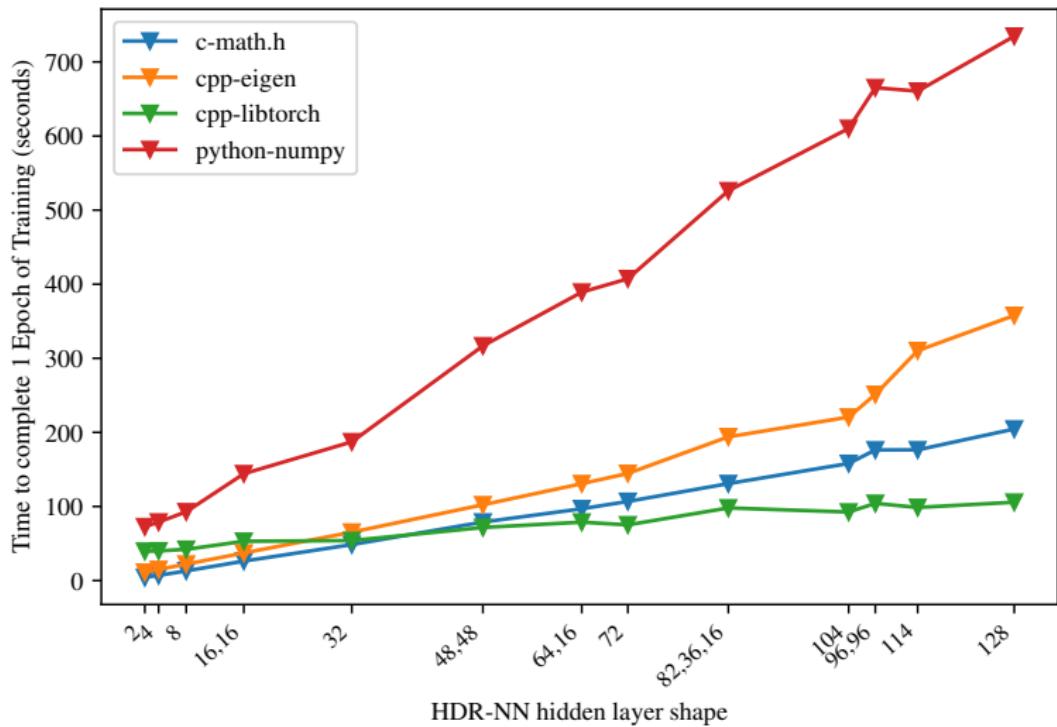
```
static void feed_forward(Network *network, float *image)
{
    //First layer is the image
    float *activations = image;

    Layer *layer = network->layers;
    while (layer != NULL)
    {
        for (int i = 0; i < layer->size; i++)
        {
            Neuron *neuron = &layer->neurons[i];
            float zvalue = 0;
            for (int j = 0; j < layer->incidents; j++)
            {
                zvalue += neuron->weights[j] * activations[j];
            }
            layer->z_values[i] = zvalue + neuron->bias;
            layer->activations[i] = sigmoid(layer->z_values[i]);
        }
        activations = layer->activations;
        layer = layer->next;
    }
}
```

Accuracy



Training Time





Development Experience

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion

- Implementing back propagation in C
- Porting libtorch to armv7hl



Future Work

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion



Reverse Engineering Efforts on the C300

NN in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Conclusion

Development Experience