



Training Neural Networks on Embedded Devices Targeting Embedded Environments

Prasanth Shaji, Deepak Venkataram

Master's Thesis
Uppsala University

August 21, 2023



Outline

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

1 Neural Networks in Embedded Devices

2 Embedded Linux

3 Neural Networks

4 Benchmark Applications

5 Discussion



UPPSALA
UNIVERSITET

Neural Network Applications on Embedded Devices

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Embedded Devices





Goal : Neural Network training on Scania ECU

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Goal : Neural Network training on Scania ECU

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the ECU?

- Repurpose existing ECU units to perform neural network related tasks
- Training on devices to reduce network bandwidth usage, ensure data privacy, etc.

Scope Repurpose Scania ECU, train neural network model, assess training performance



Goal : Neural Network training on Scania ECU

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the ECU?

- Repurpose existing ECU units to perform neural network related tasks
- Training on devices to reduce network bandwidth usage, ensure data privacy, etc.

Scope Repurpose Scania ECU, train neural network model, assess training performance



Goal : Neural Network training on Scania ECU

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the ECU?

- Repurpose existing ECU units to perform neural network related tasks
- Training on devices to reduce network bandwidth usage, ensure data privacy, etc.

Scope Repurpose Scania ECU, train neural network model, assess training performance



ARM Boards : Examples

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

SoC	Application Domain	Architecture	Processor Core
nRF51822	Ultra low power, BLE	ARM v6-M	ARM Cortex M0
AM2732	Automotive	ARM v7-R	ARM Cortex R5
AM3358	Industrial / IoT	ARM v7-A	ARM Cortex A8
i.MX6S	Multimedia applications	ARM v7-A	ARM Cortex A9
BCM2711	Educational / IoT	ARM v8-A	ARM Cortex-A72
Snapdragon - Kryo 240	Smartphones	ARM v8-A	ARM Cortex-A73, ARM Cortex-A53
Apple A16 Bionic	iPhone 14 Pro	ARM v8.6-A	APL1W10
Apple M1	Macbooks / Mac Mini / iPads	ARM v8.5-A	APL1102



ARM Boards : Outline

ARM | Silicon Vendor | System Maker

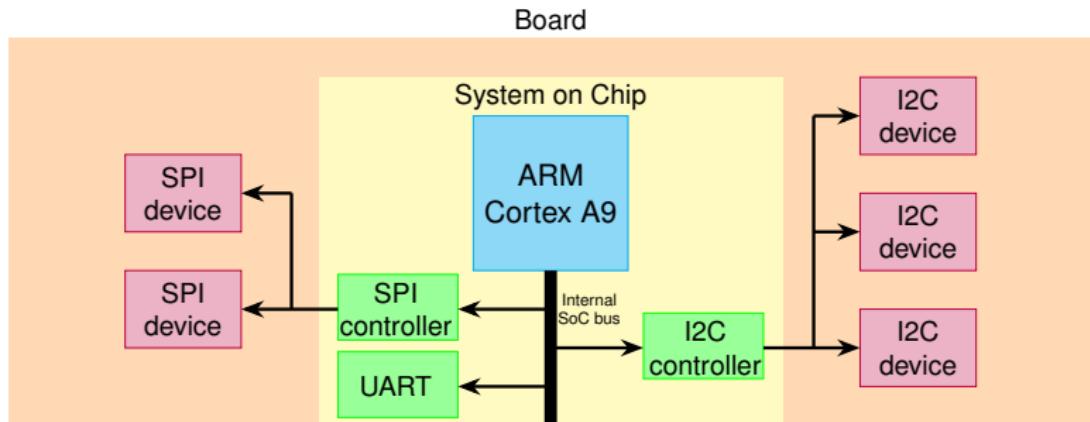
Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Embedded Linux

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Kernel, Bootloaders, Device Trees
Cross Compilation, Toolchains, Root Filesystems

Operating System for Embedded Devices
- Customized for Specific Hardware



Embedded Linux

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Kernel, Bootloaders, Device Trees
Cross Compilation, Toolchains, Root Filesystems

Operating System for Embedded Devices
- Customized for Specific Hardware



Build Systems

Neural
Networks in
Embedded
Devices

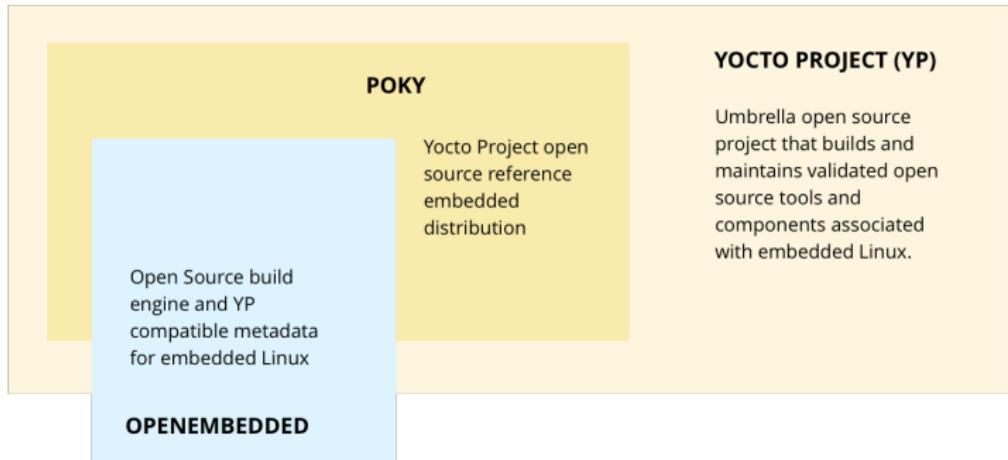
Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Yocto Project





The Yocto Project

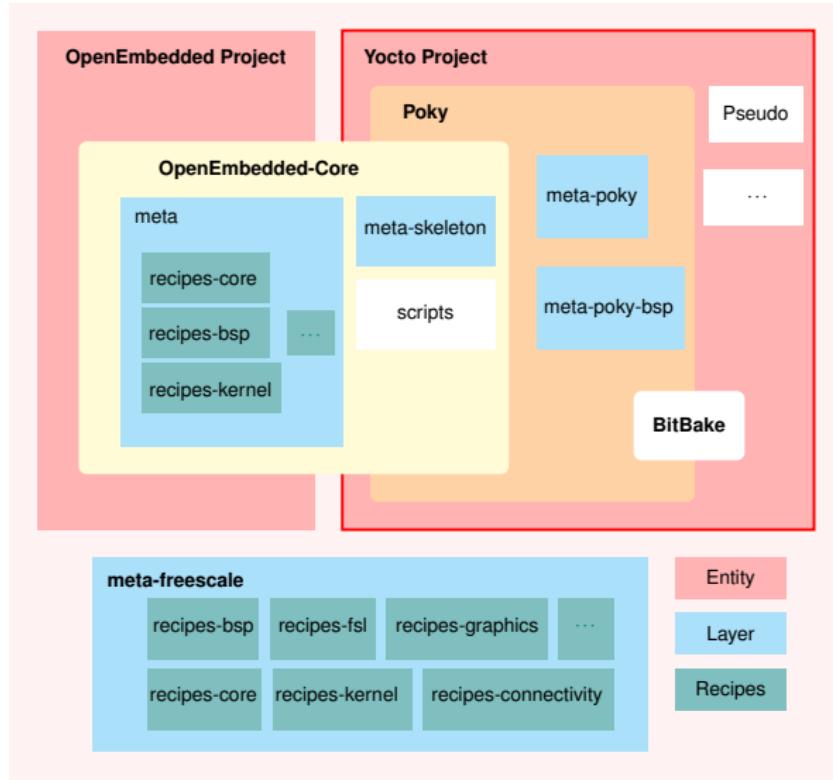
Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

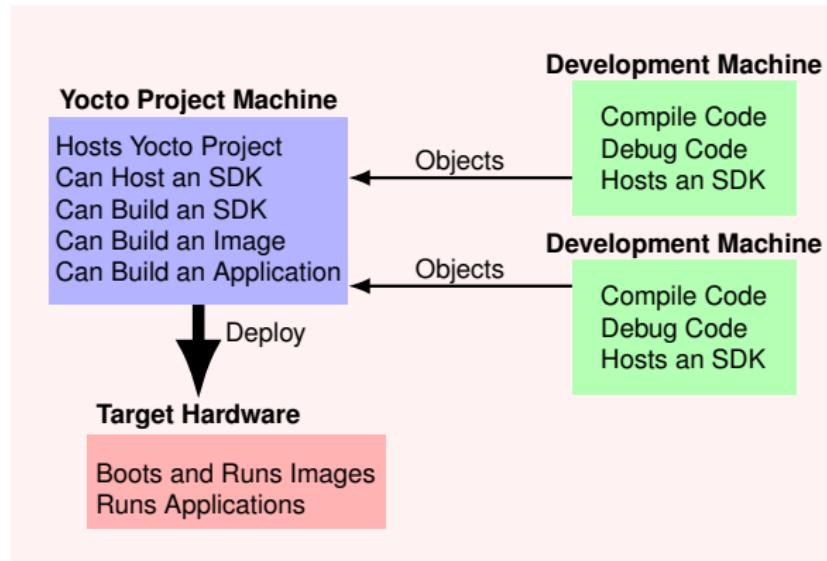
Benchmark
Applications

Discussion





Yocto Project : Development





Software Versions

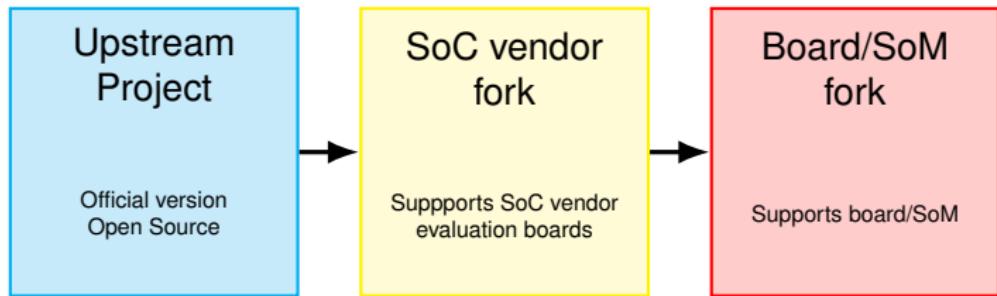
Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion



Scania ECU to MCIMX6Q-SDB

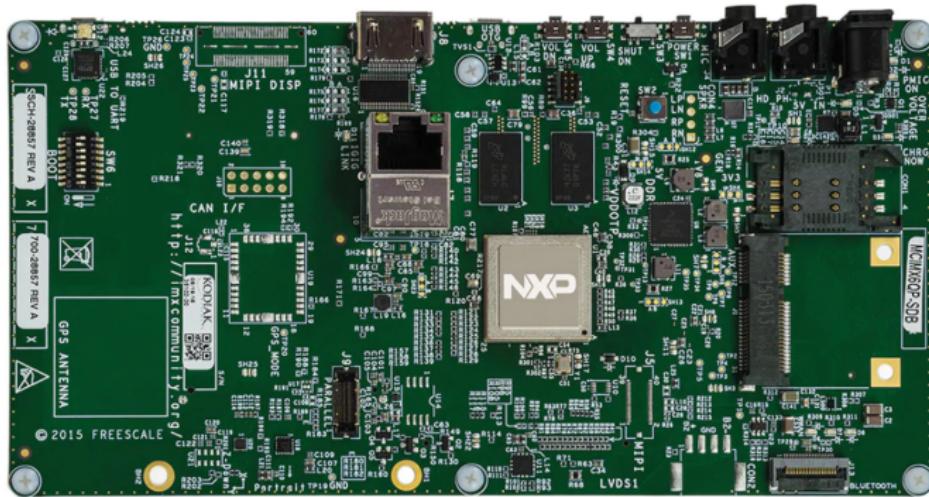
Neural Networks in Embedded Devices

Embedded Linux

Neural Networks

Benchmark Applications

Discussion





Handwritten Digit Recognition

Neural Network



MNIST



Neural Network Training

Neural
Networks in
Embedded
Devices

Embedded
Linux

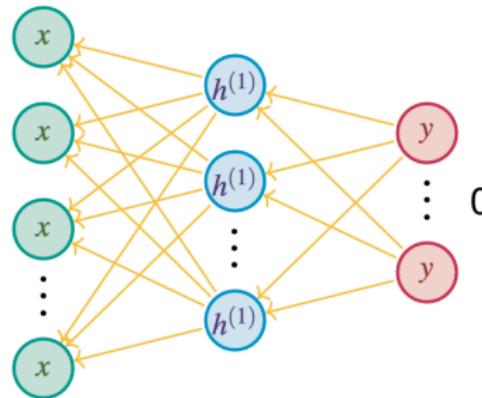
Neural
Networks

Benchmark
Applications

Discussion



label = 0





Learning Algorithm

Algorithm 1 Mini Batch Gradient Descent with learning rate η and the Mean Squared Error (MSE) cost function

Require: initial weights $w^{(0)}$, number of epochs E , batch size B , training data with T entries

Ensure: final weights $w^{(E*T)}$

```
for  $e = 0 \rightarrow E - 1$  do
    for  $b = 0 \rightarrow T/B$  do
        for  $t = b * B \rightarrow (b + 1) * B$  do
            estimate  $\nabla \mathcal{L}(w^{(t)})$ 
            compute  $\Delta w^{(b)} += -\nabla \mathcal{L}(w^{(t)})$ 
        end for
         $w^{(e+1)} := w^{(e)} + \eta \Delta w^{(e)}$ 
    end for
end for
return  $w^{(T)}$ 
```

▷ \mathcal{L} here is MSE



Neural Network Inference

Neural
Networks in
Embedded
Devices

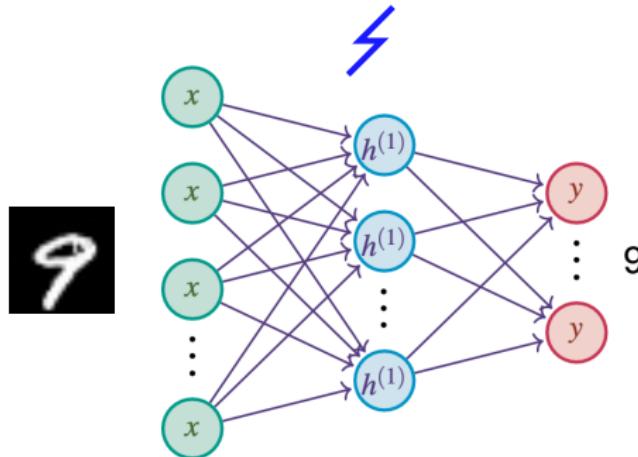
Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

9





Traditional Paradigm

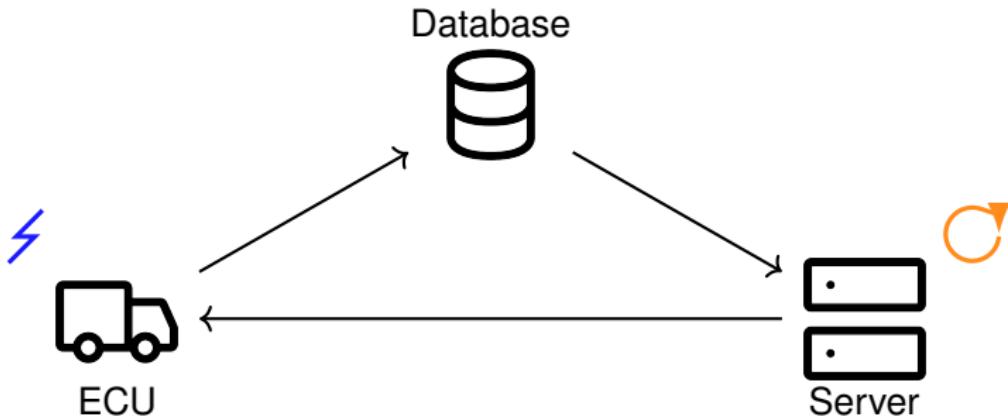
Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Federated Learning

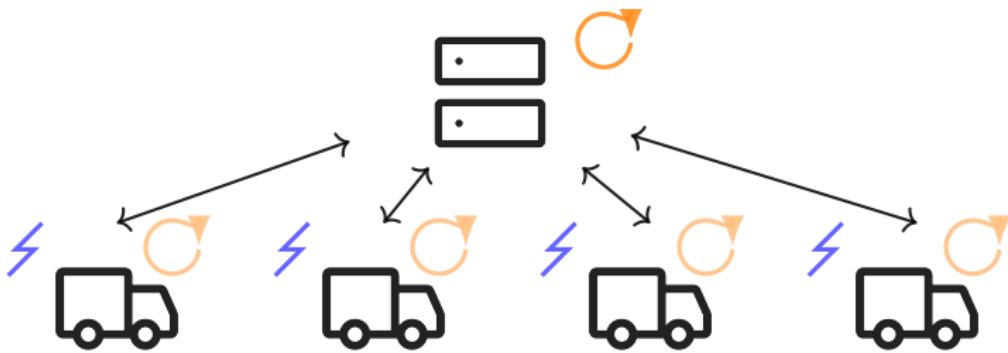
Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





C based HDR-NN

github.com/paperbound/hdrnn-arm

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

```
/*-HDR-Neural-Network-*/
typedef struct
{
    float *bias;
    float **weights;
    float **nabla_w;
} Neuron;

typedef struct LayerT
{
    int size;
    int incidents;
    Neuron **neurons;
    float **activations;
    float **z_values;
    float **nabla_b;
    struct LayerT *next;
    struct LayerT *previous;
} Layer; // Network.layers.except.for.input

typedef struct
{
    Layer **layers;
    int depth;
} Network; // HDRNN
```

```
static void feed_forward(Network *network, float *image)
{
    // First layer is the image
    float *activations = image;

    Layer *layer = network->layers;
    while (layer != NULL)
    {
        for (int i = 0; i < layer->size; i++)
        {
            Neuron *neuron = &layer->neurons[i];
            float zvalue = 0;
            for (int j = 0; j < layer->incidents; j++)
            {
                zvalue += neuron->weights[j] * activations[j];
            }
            layer->z_values[i] = zvalue + neuron->bias;
            layer->activations[i] = sigmoid(layer->z_values[i]);
        }
        activations = layer->activations;
        layer = layer->next;
    }
}
```



HDR-NN configurability

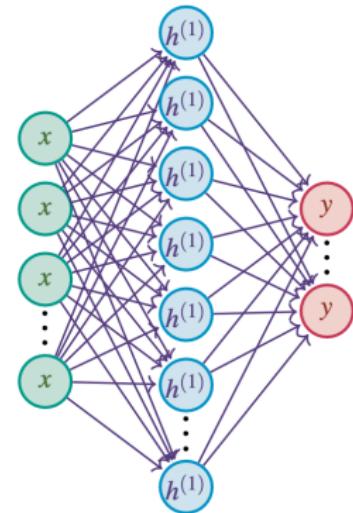
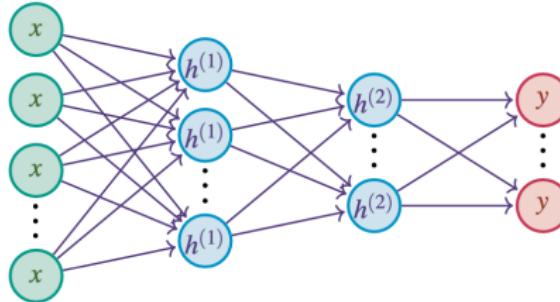
Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





HDR-NN configurability

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

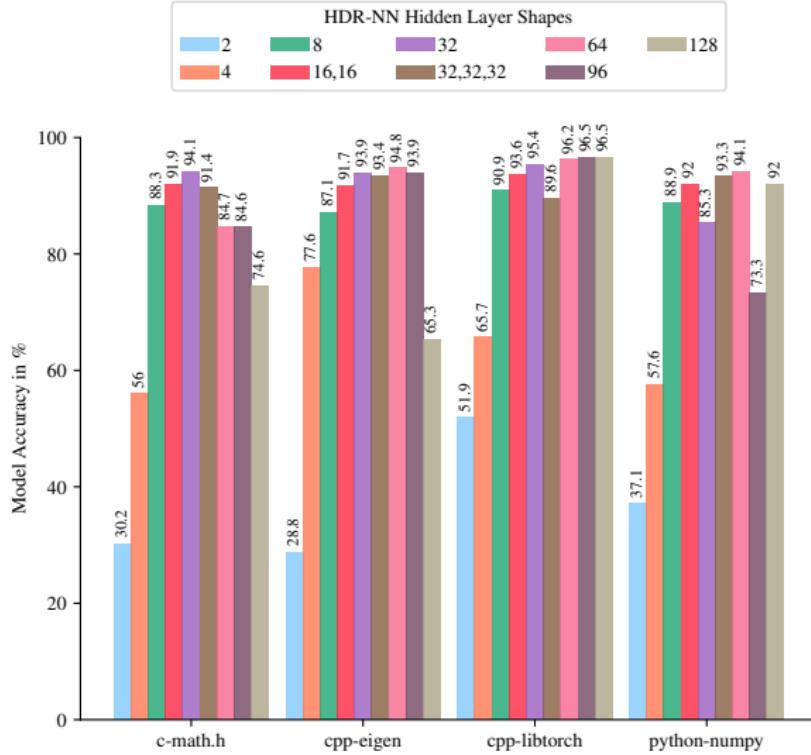
```
→ ./bin/hdrnn
nothing to do
Usage: ./bin/hdrnn

      <command> [<args>]

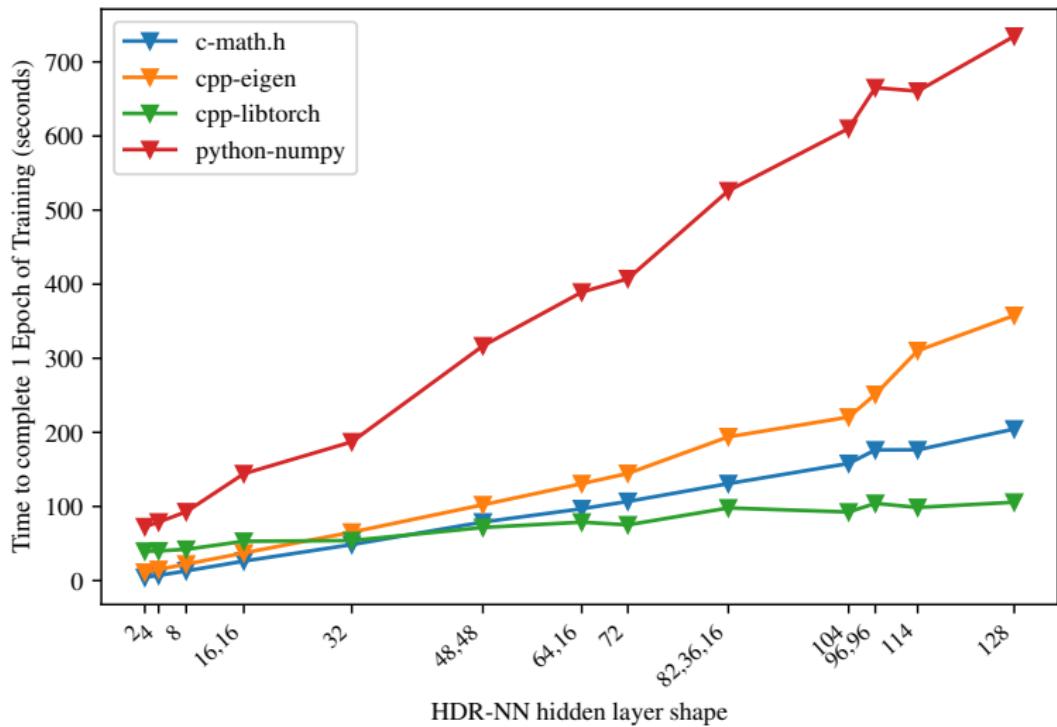
Commands:
  infer [-i, --image IMAGE_PATH] [-n, --net c-math.nn]
  train [-s, --shape 32] [-e, --epochs 30]
        [-q, --quiet] [-bs, --batch_size=10]
        [-lr, --learning_rate 3] [-n, --net c-math.nn]
```



Accuracy



Training Time





Development Journey

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- Linux, software packages for MCIMX6Q
- Implementing back propagation in C
- Implementing HDR-NN in C, C++ / Eigen
- Compiling libtorch for armv7hl



Software Development on Embedded Devices

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- Embedded Platforms are fragmented in both *hardware* and *software*
- Challenging to advance software
- Research Approaches : Tiny ML, Tiny Training Engine



Future Work

Neural
Networks in
Embedded
Devices

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- Reverse Engineering Efforts on the ECU
- Performance engineering the C based implementation
- Adding more target devices
- Adding more implementations (TTE)