



Training Neural Networks on Embedded Devices Targeting Embedded Environments

Prasanth Shaji, Deepak Venkataram

Master's Thesis
Uppsala University



Outline

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

1 Introduction

2 Embedded Linux

3 Neural Networks

4 Benchmark Applications

5 Discussion



Embedded Devices



Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion



Neural Network Applications on Embedded Devices

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Goal : Neural Network Training on Scania ECU



Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion



Traditional Paradigm

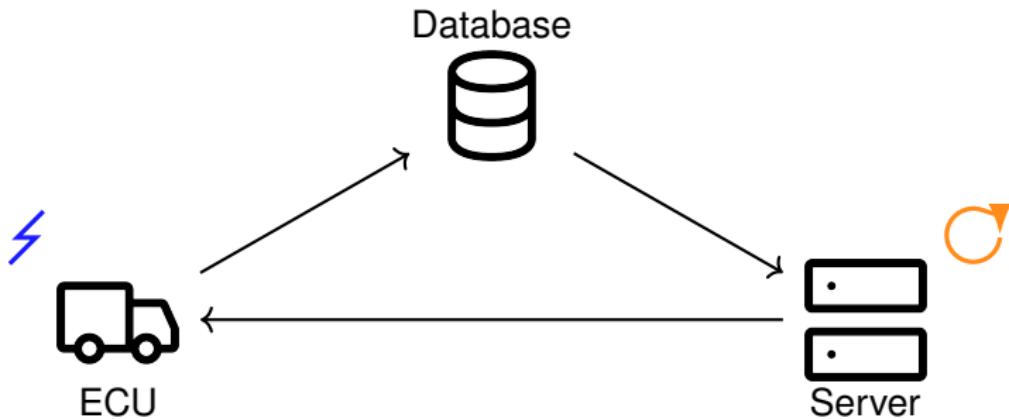
Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Federated Learning

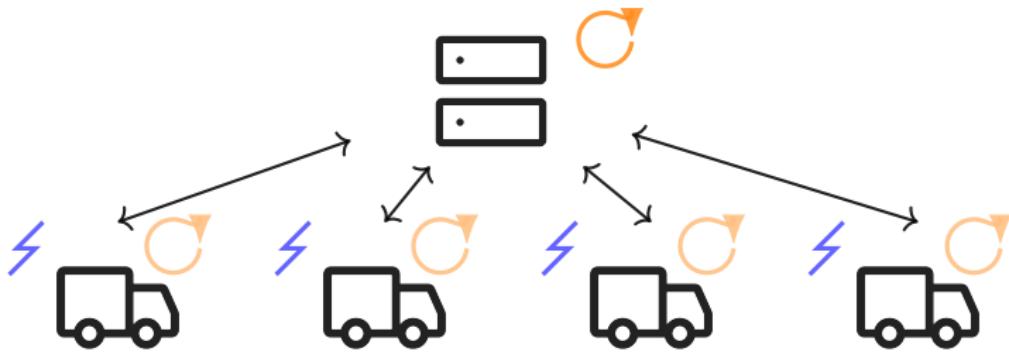
Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Goal : Neural Network Training on Scania ECU

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the ECU?

- Training on devices reduces network usage, ensures data privacy, etc.
- Repurpose existing ECUs to perform neural network related tasks

Scope Repurpose Scania ECU, train neural network model, assess training performance



Goal : Neural Network Training on Scania ECU

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the ECU?

- Training on devices reduces network usage, ensures data privacy, etc.
- Repurpose existing ECUs to perform neural network related tasks

Scope Repurpose Scania ECU, train neural network model, assess training performance



Goal : Neural Network Training on Scania ECU

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the ECU?

- Training on devices reduces network usage, ensures data privacy, etc.
- Repurpose existing ECUs to perform neural network related tasks

Scope Repurpose Scania ECU, train neural network model, assess training performance



Goal : Neural Network Training on Scania ECU

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Predictive vehicle maintenance using data from a fleet of connected trucks

Why Training on the ECU?

- Training on devices reduces network usage, ensures data privacy, etc.
- Repurpose existing ECUs to perform neural network related tasks

Scope *Repurpose Scania ECU, train neural network model, assess training performance*



Embedded Devices : Hardware Outline

ARM | Silicon Vendor | System Maker

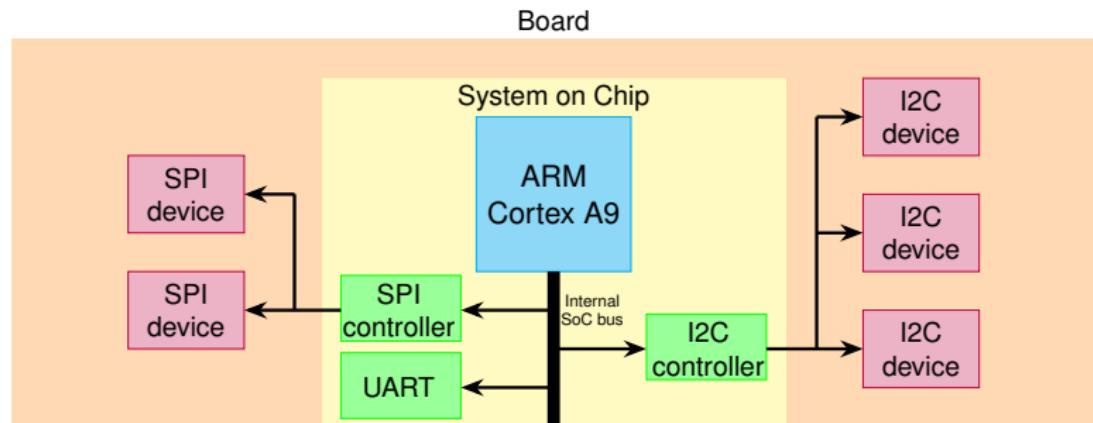
Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





ARM Boards : Examples

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

SoC	Application Domain	Architecture	Processor Core
nRF51822	Ultra low power, BLE	ARM v6-M	ARM Cortex M0
AM2732	Automotive	ARM v7-R	ARM Cortex R5
AM3358	Industrial / IoT	ARM v7-A	ARM Cortex A8
i.MX6S	Multimedia applications	ARM v7-A	ARM Cortex A9
BCM2711	Educational / IoT	ARM v8-A	ARM Cortex-A72
Snapdragon 460	Smartphones	ARM v8-A	ARM Cortex-A73, A53
Google Tensor	Pixel 6	ARM v8.2-A	ARM Cortex-X1, A76, A55
Apple A16 Bionic	iPhone 14 Pro	ARM v8.6-A	APL1W10
Apple M1	Macbooks / Mac Mini / iPads	ARM v8.5-A	APL1102



Embedded Linux

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Operating System for Embedded Devices
- Customized for Specific Hardware

Kernel, Bootloaders, Device Trees
Cross Compiling Toolchains, Root Filesystems



Embedded Linux

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Operating System for Embedded Devices
- Customized for Specific Hardware

Kernel, Bootloaders, Device Trees
Cross Compiling Toolchains, Root Filesystems



Build Systems

Introduction

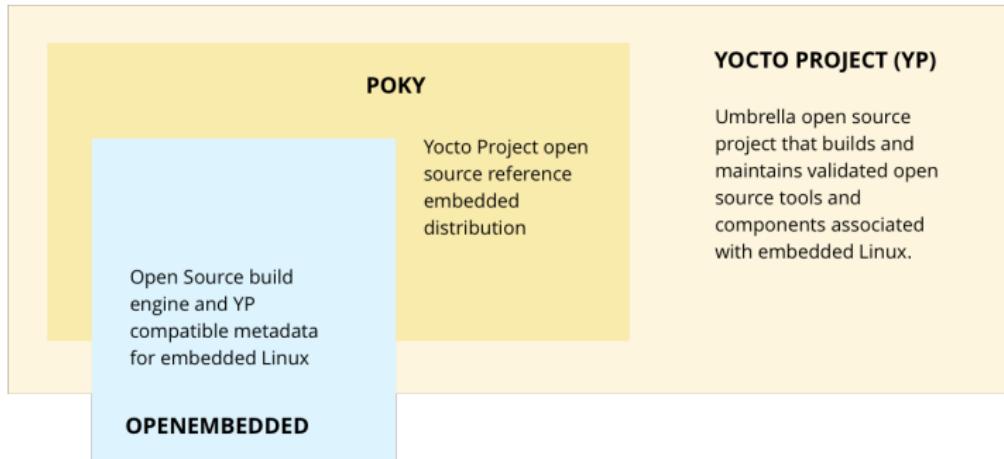
Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Yocto Project





Development Setup

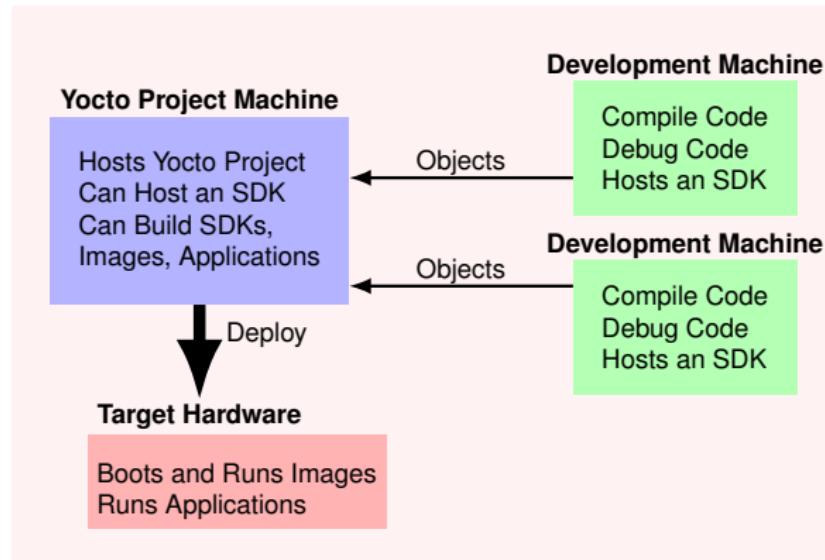
Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Software Versions

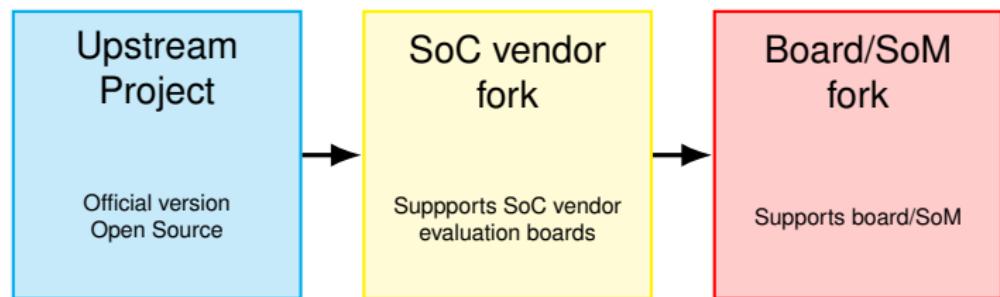
Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Scania ECU to MCIMX6Q-SDB

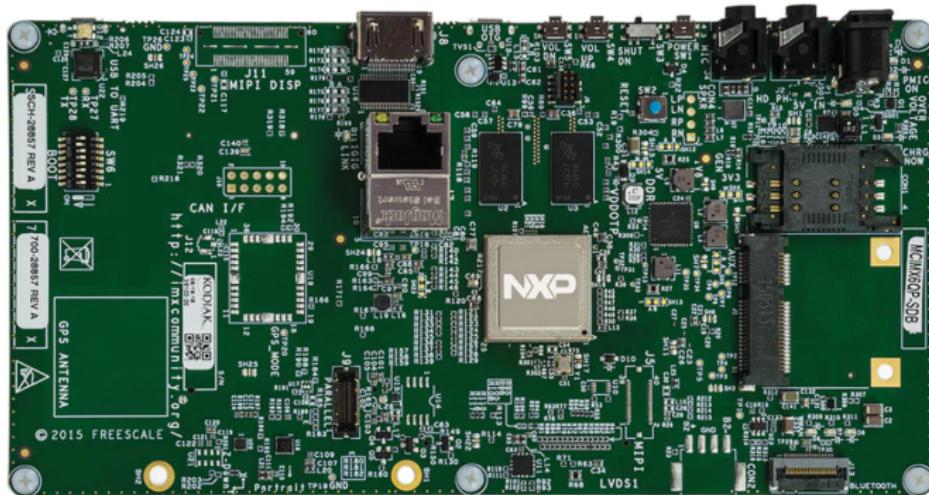
Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





Handwritten Digit Recognition

Neural Network

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion



MNIST



Neural Network

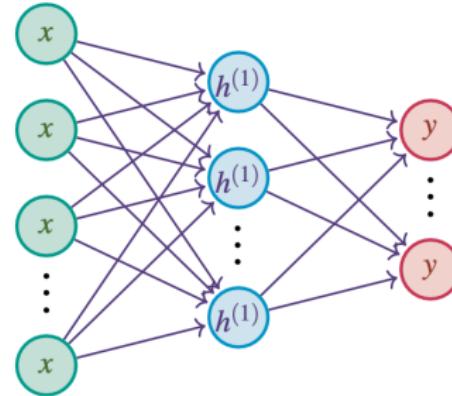
Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion



$$y_k = \sigma \left(\sum_{j=0}^m w_{kj} x_j \right)$$



Neural Network Training

Introduction

Embedded
Linux

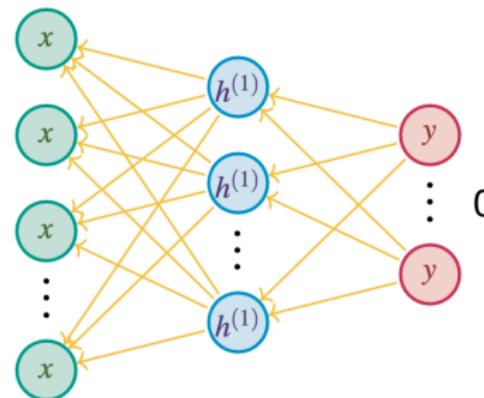
Neural
Networks

Benchmark
Applications

Discussion



label = 0





Learning Algorithm

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Algorithm 1 Mini Batch Gradient Descent with learning rate η and the Mean Squared Error (MSE) cost function

Require: initial weights $w^{(0)}$, number of epochs E , batch size B , training data with T entries

Ensure: final weights $w^{(E*T)}$

```
for  $e = 0 \rightarrow E - 1$  do
    for  $b = 0 \rightarrow T/B$  do
        for  $t = b * B \rightarrow (b + 1) * B$  do
            estimate  $\nabla \mathcal{L}(w^{(t)})$ 
            compute  $\Delta w^{(b)} += -\nabla \mathcal{L}(w^{(t)})$ 
        end for
         $w^{(e+1)} := w^{(e)} + \eta \Delta w^{(e)}$ 
    end for
end for
return  $w^{(T)}$ 
```

▷ \mathcal{L} here is MSE



Neural Network Inference

Introduction

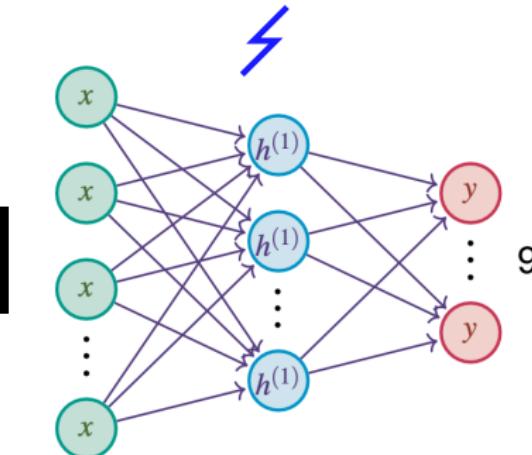
Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

9





HDR-NN Implementations

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

github.com/paperbound/hdrnn-arm

- C based HDR-NN
- C++ / Eigen
- C++ / libtorch (PyTorch)
- Python / Numpy



C based HDR-NN

Introduction

Embedded
LinuxNeural
NetworksBenchmark
Applications

Discussion

```
/*-HDR-Neural-Network*/
typedef struct
{
    float *bias;
    float *weights;
    float *nabla_w;
} Neuron;

typedef struct LayerT
{
    int size;
    int incidents;
    Neuron *neurons;
    float *activations;
    float *z_values;
    float *nabla_b;
    struct LayerT *next;
    struct LayerT *previous;
} Layer; // Network.layers except for input

typedef struct
{
    Layer *layers;
    int depth;
} Network; // HDRNN
```

```
static void feed_forward(Network *network, float *image)
{
    // First layer is the image
    float *activations = image;

    Layer *layer = network->layers;
    while (layer != NULL)
    {
        for (int i = 0; i < layer->size; i++)
        {
            Neuron *neuron = &layer->neurons[i];
            float zvalue = 0;
            for (int j = 0; j < layer->incidents; j++)
            {
                zvalue += neuron->weights[j] * activations[j];
            }
            layer->z_values[i] = zvalue + neuron->bias;
            layer->activations[i] = sigmoid(layer->z_values[i]);
        }
        activations = layer->activations;
        layer = layer->next;
    }
}
```



C++ / Eigen based HDR-NN

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- Eigen is a C++ template library for linear algebra
- Analogous to Python / Numpy



C++ / libtorch based HDR-NN

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- PyTorch doesn't support MCIMX6Q
- Porting libtorch using QEMU



HDR-NN configurability

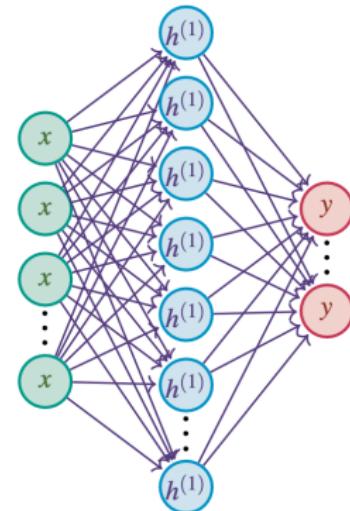
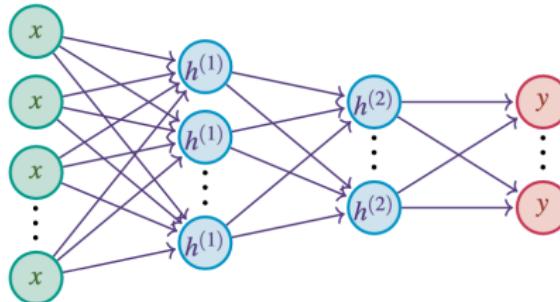
Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion





HDR-NN configurability

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

```
→ ./bin/hdrnn
nothing to do
Usage: ./bin/hdrnn

      <command> [<args>]

Commands:
  infer [-i, --image IMAGE_PATH] [-n, --net c-math.nn]
  train [-s, --shape 32] [-e, --epochs 30]
        [-q, --quiet] [-bs, --batch_size=10]
        [-lr, --learning_rate 3] [-n, --net c-math.nn]
```



Experiment Considerations

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- Try multiple shapes m
- Compare Accuracies
- Compare Execution Time



Accuracy

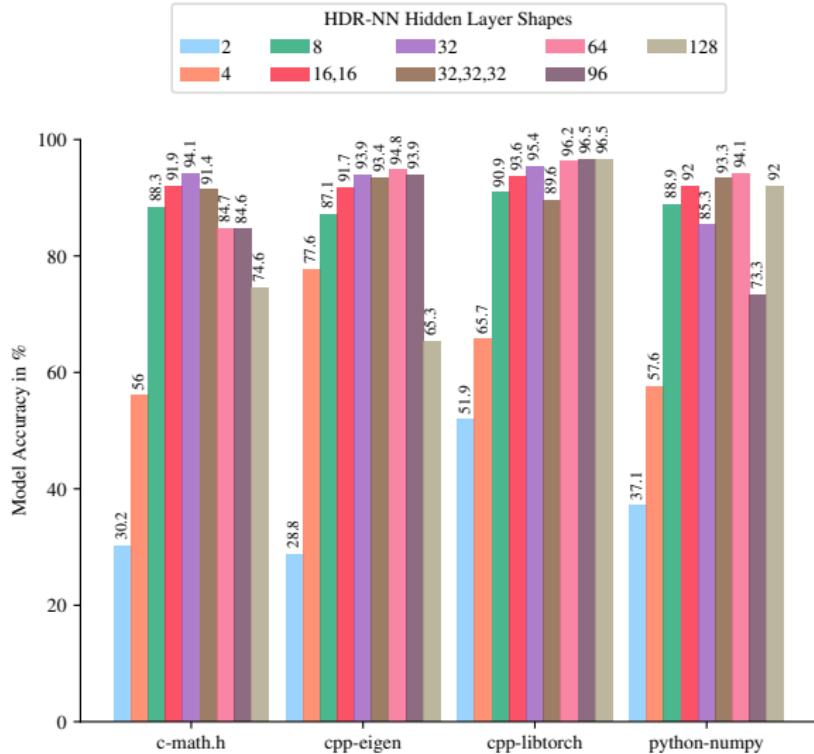
Introduction

Embedded Linux

Neural Networks

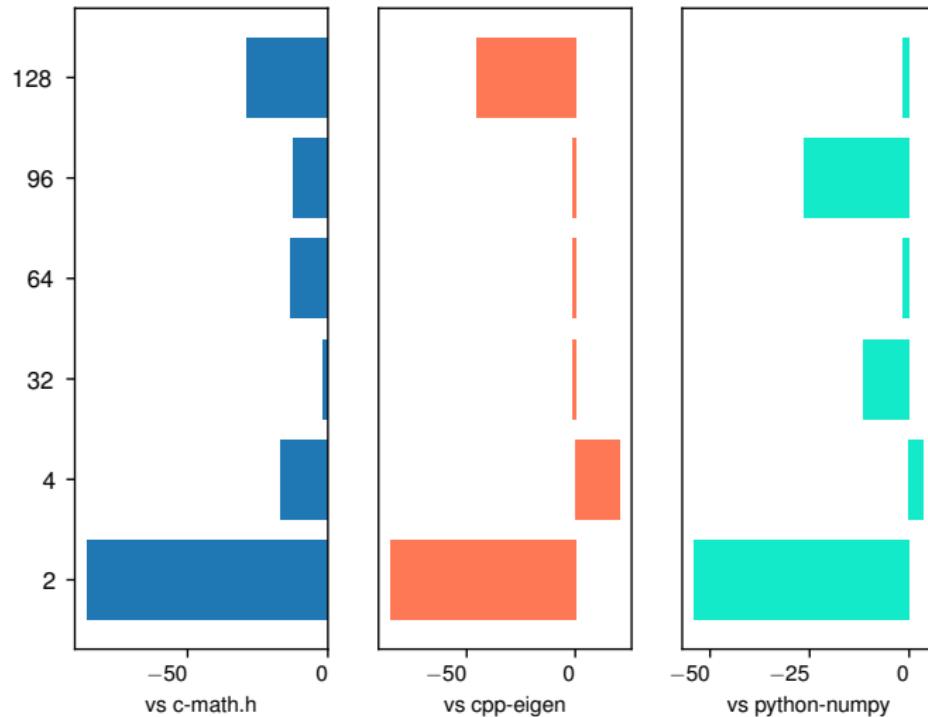
Benchmark Applications

Discussion



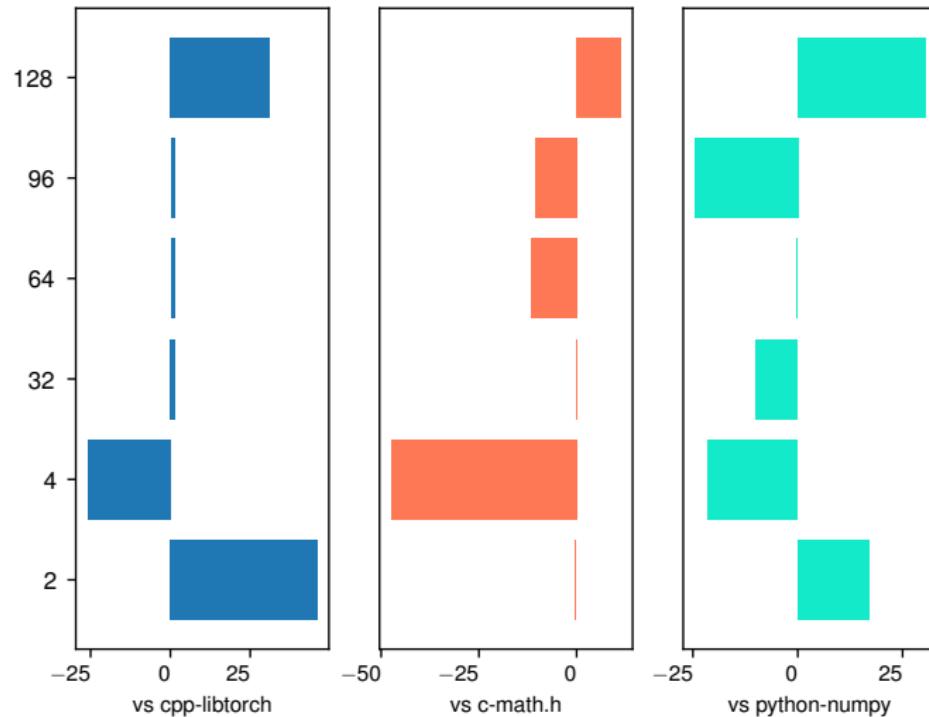


Comparing Accuracies of cpp-libtorch



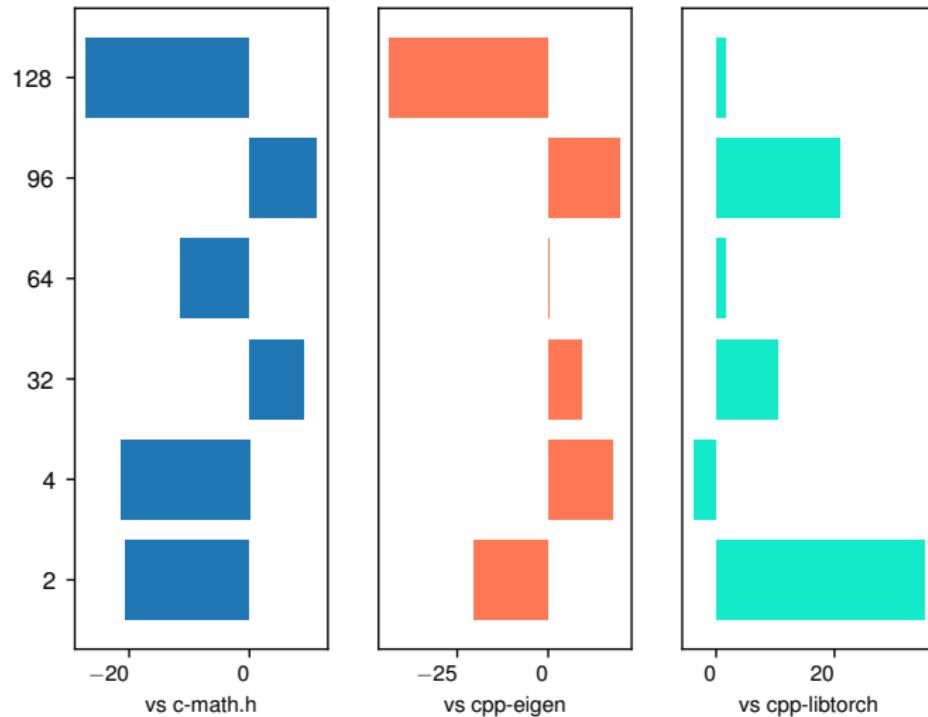


Comparing Accuracies of cpp-eigen



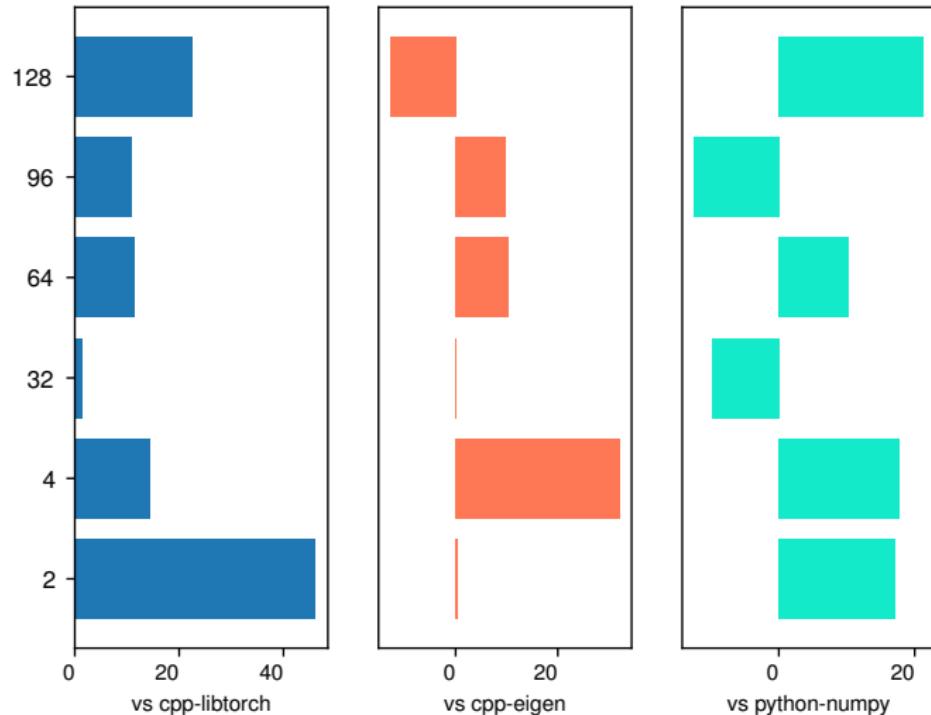


Comparing Accuracies of python-numpy





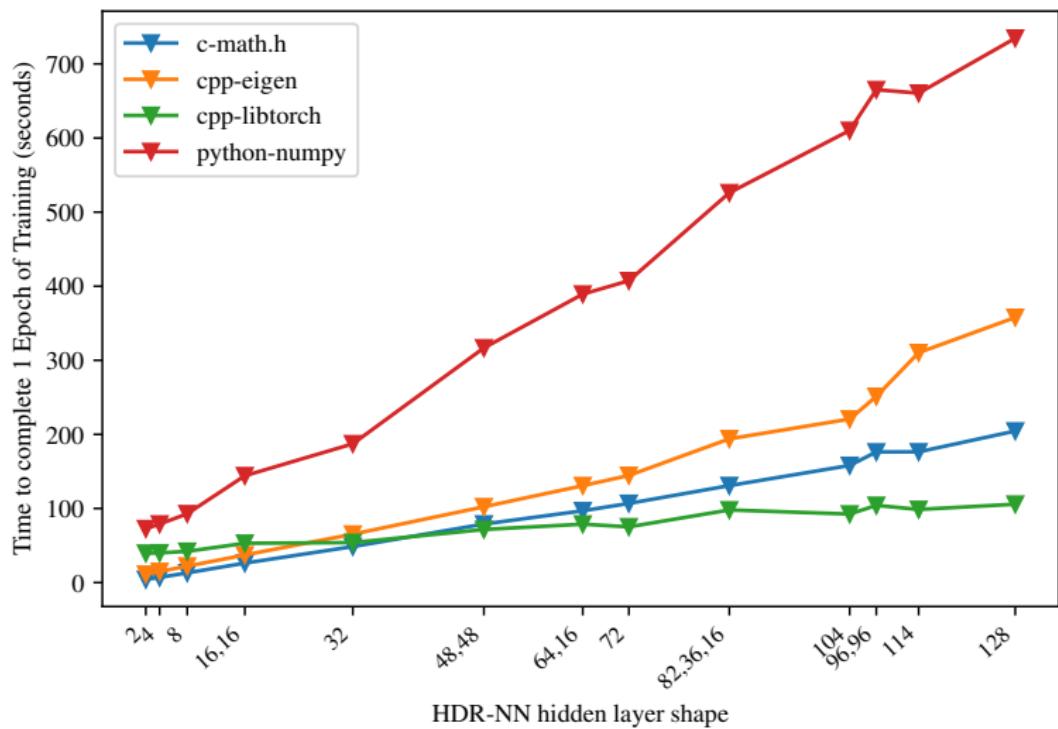
Comparing Accuracies of c-math.h



Training Time

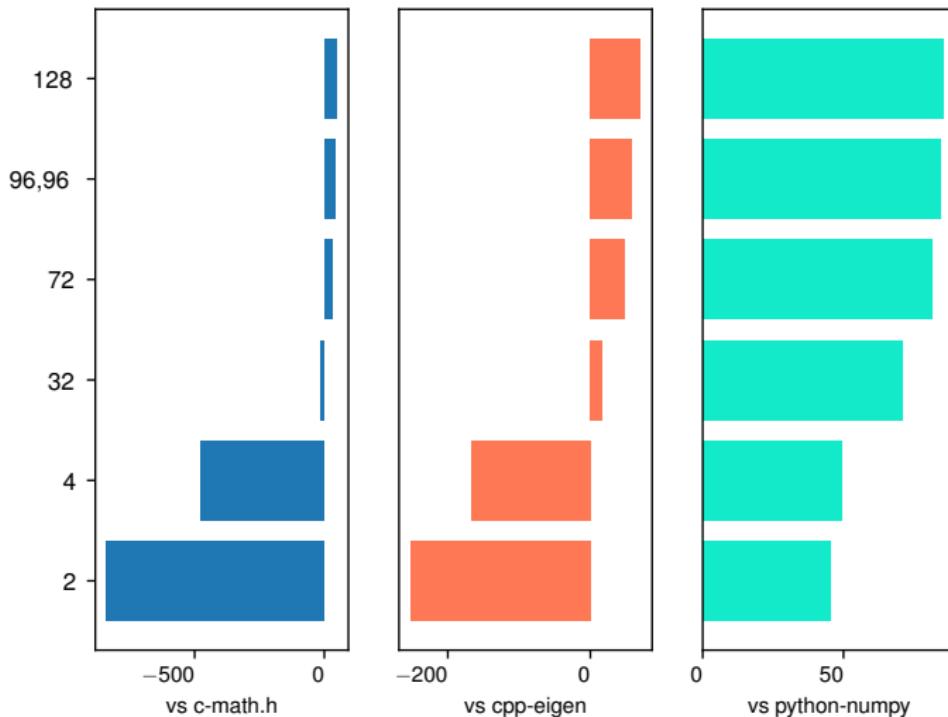
Introduction
Embedded Linux
Neural Networks
Benchmark Applications

Discussion



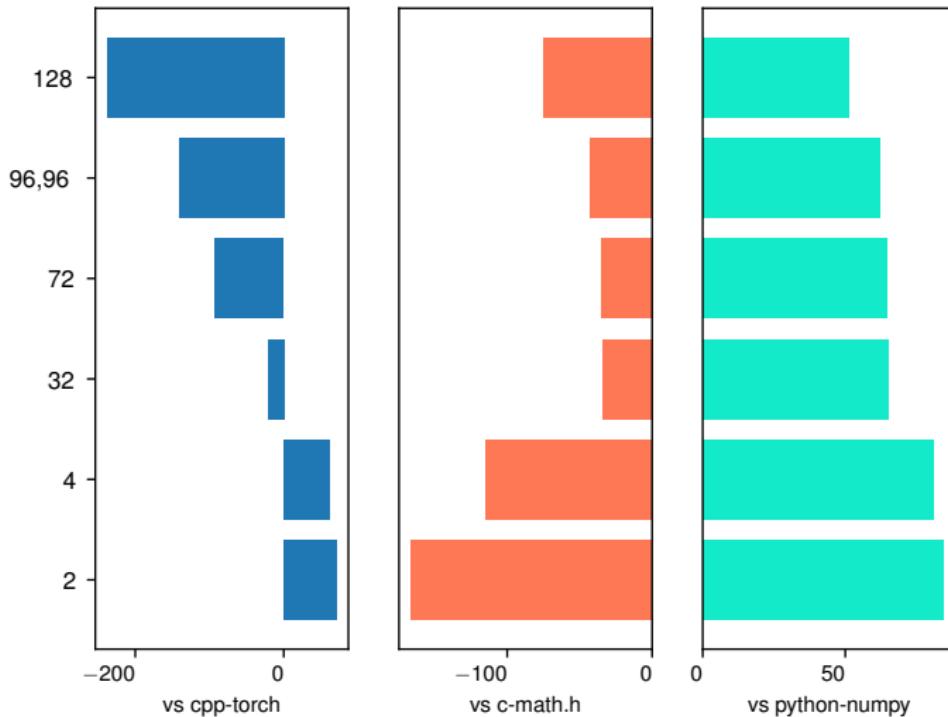


Comparing Execution Times of cpp-libtorch



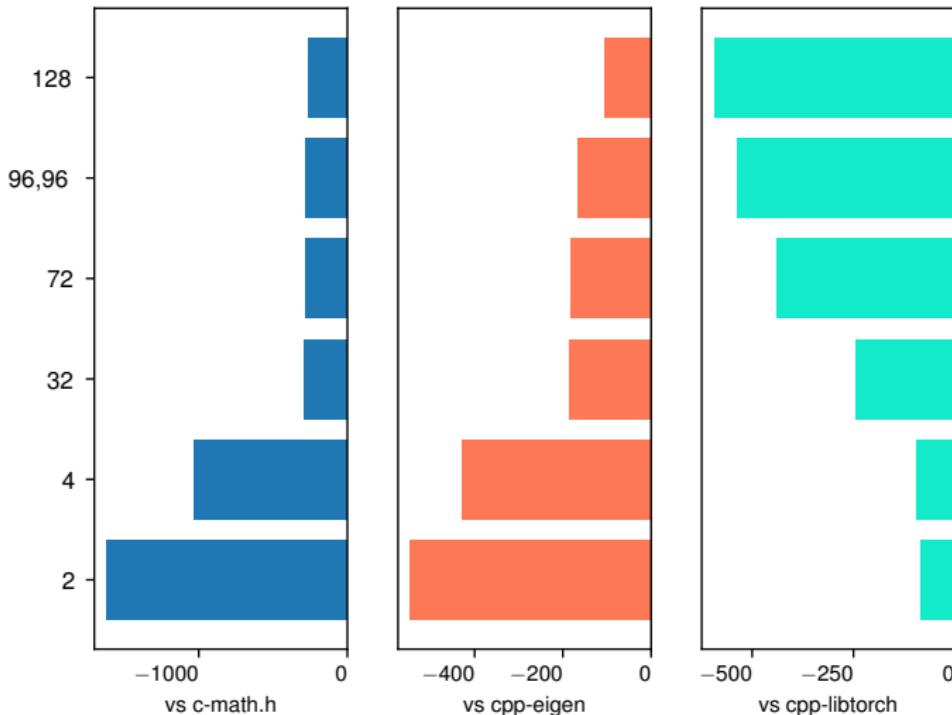


Comparing Execution Times of cpp-eigen



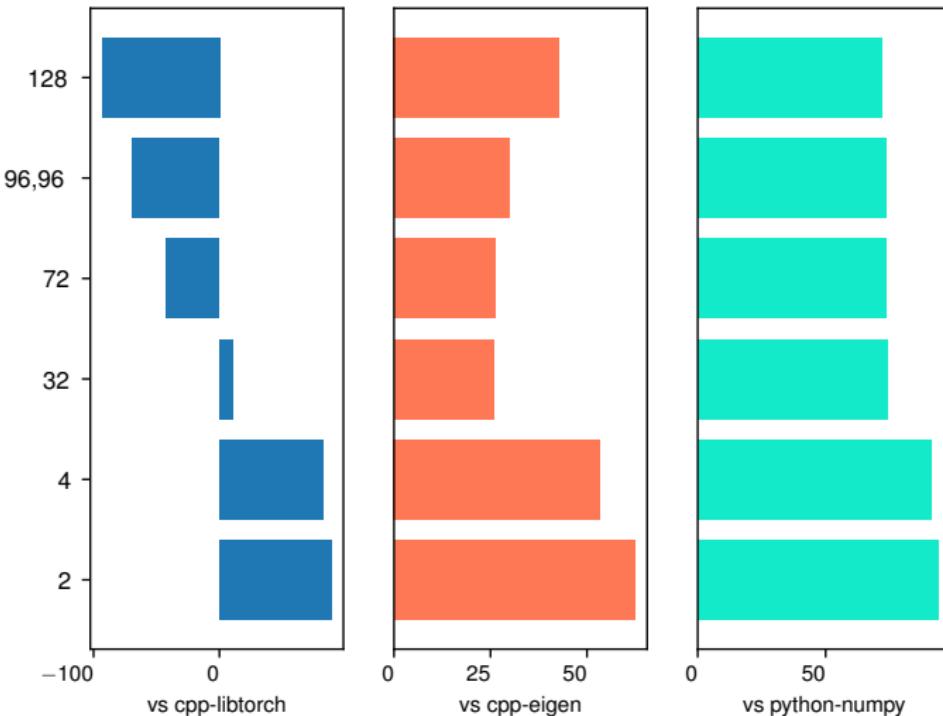


Comparing Execution Times of python-numpy





Comparing Execution Times of c-math.h





Reverse Engineering Scania ECU

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- Goal : Get a vanilla U-Boot running on ECU
- Extract device tree information
- Silent bootloader, No visibility into Board/SoM fork



Neural Network Software Development on Embedded Devices

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- Embedded Platforms are fragmented in both *hardware* and *software*
- C is easier to target harder to Implement
- Python commonly found in Neural Network Research, harder to port



Future Work

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

- Reverse Engineering Efforts on the ECU
- Performance engineering the C based implementation
- Adding more target devices
- Adding more implementations (Tensorflow, TTE)



End

Introduction

Embedded
Linux

Neural
Networks

Benchmark
Applications

Discussion

Thank you
Questions?